Closed-Form Moment Reduction of 3D Space-Charge Fields for Fast Feedback in High-Intensity Hadron Rings

Hyo Eun Ahn*

November 3, 2026

Abstract

In high-intensity hadron accumulator rings, rapid-cycling synchrotrons, and low-to-medium energy injector boosters—such as the Los Alamos Proton Storage Ring (PSR), the Spallation Neutron Source (SNS) accumulator ring, the CERN Proton Synchrotron Booster (PSB), and FAIR SIS18—intense self-induced Coulomb repulsion inside dense particle bunches ($\gamma \sim 1.15$–3.1) drives large incoherent space-charge tune shifts ($\Delta Q_{\mathrm{sc}} \sim -0.1$ to -0.4), resulting in envelope mismatch resonances, transverse emittance growth, and halo generation. While active damping of dipolar centroid oscillations has been standard practice since Simon van der Meer's stochastic cooling, real-time active compensation of higher-order space-charge fields remains computationally challenging due to the millisecond-to-second latency of direct $O(N^2)$ Coulomb summation and grid-based $O(N_p + G \log G)$ Particle-in-Cell (PIC) solvers relative to the sub-microsecond to microsecond revolution period ($T_{\mathrm{rev}} \approx 0.36$–$1.0\,\mu$s, or $\leq 5\,\mu$s for a 5-turn envelope control window). In this paper, building on the author's 1994 Los Alamos space-charge moment formulations, we establish a closed-form reduced-order moment framework that evaluates 3D interior space-charge potential and electric-field polynomial coefficients via 35×35 and 34×35 matrix-vector kernels (~ 1.6–$5.8\,\mu$s compute latency on a single CPU core, ~ 5.4–$20\,\mu$s including Python-level Zero-Self-Force indexing, and $< 1.0\,\mu$s projected on FPGA hardware for $1,247$ multiply-accumulate operations). By expanding the bunch distribution into $M = 35$ spatial moments (all 3D monomials up to 4th order) and regularizing the underdetermined N-point mesh inversion ($\mathrm{rank}(\boldsymbol{\Omega}) = 35$) with a normalized physical reference prior ($\sum_{i=1}^{N} \mu_i = 1, \mu_i > 0$), strict pointwise positivity ($\boldsymbol{a} \equiv \boldsymbol{\mu} > \mathbf{0}$) is guaranteed at the reference equilibrium and fringe negative-charge overshoot is suppressed by $3.04\times$ ($0.2669 \to 0.0876$) in our 3D anisotropic bunch benchmark compared to Channell's uniform prior. Using reference-prior linearization ($\boldsymbol{w} = \boldsymbol{\mu}$), we demonstrate the exact algebraic "fall-out" of the mesh dimension N, collapsing the prior-weighted polynomial projection of the Plummer-softened Coulomb kernel into precomputed 35×35 potential ($\boldsymbol{v} = \boldsymbol{K}\boldsymbol{s} + \boldsymbol{v}_0$) and 34×35 electric-field ($\boldsymbol{v}_E^{\mathrm{prior}} = \boldsymbol{K}_E\boldsymbol{s} + \boldsymbol{v}_{E,0}$) matrices. Furthermore, because the uncorrected prior-weighted operator guarantees $\langle \boldsymbol{E}_{\mathrm{approx}} \rangle_{\boldsymbol{\mu}} = \langle \boldsymbol{E}_{\mathrm{exact}} \rangle_{\boldsymbol{\mu}}$ rather than vanishing net force on a perturbed state $\boldsymbol{a} = \boldsymbol{\mu} + \delta\boldsymbol{a}$ (where $\boldsymbol{\mu}^T \boldsymbol{G} \delta\boldsymbol{a} \neq \mathbf{0}$), we introduce an exact $O(M_E)$ moment-space Zero-Self-Force (ZSF) dipolar offset projection that restores $\langle \boldsymbol{E}_{\mathrm{approx}}^{\mathrm{ZSF}} \rangle_{\boldsymbol{a}} = \mathbf{0}$ to machine precision (4.02×10^{-16}) without accessing the N-point mesh. In an idealized zero-latency multi-turn 2D transverse Leapfrog tracking and de-confounded ten-mode ablation study driven by self-consistent $O(N_p^2)$ 3D pairwise Plummer-Coulomb forces from identical initial conditions, we separate the physical roles of centroid-frame alignment and ZSF dipolar projection: aligning the polynomial evaluation frame with the moving beam centroid suppresses off-axis multipole feed-down (reducing RMS emittance growth from +40.65% open-loop to +17.29% for 11-term multipole correctors, +13.45% for the 34-basis operator, and +6.40% with a 2nd-moment rate damper), while Corollary 4's ZSF projection eliminates both polynomial and applied net self-force ($\sim 10^{-16}$) and suppresses spurious center-of-mass trajectory distortion (0.049–$0.081 \to \sim 10^{-17}$) purely from the 35 moments without altering central internal moments. The physical field accuracy of the 4th-order polynomial basis is $R^2 = 0.8878$ ($\epsilon_E = 33.50\%$ RMS, ZSF operator, bunch core

*Email: hyo.eun.ahn@proton.me

$r \leq 0.45$), inherent to the fixed-prior degree-$P \leq 4$ truncation; all multi-turn emittance figures are from an idealized zero-latency benchmark and do not constitute a full hardware demonstration.

Keywords: Space-Charge Field Reduction, High-Intensity Hadron Rings, Injector Boosters, Moment Approach, Reduced-Order Modeling, Fast Multipole Feedback.

—

1 Introduction

The ultimate beam brightness (N_b/ϵ_n) delivered by high-power proton drivers and hadron collider injector chains is fundamentally limited by collective space-charge forces at low and intermediate energies [1, 2]. In the laboratory frame, the self-generated magnetic force of a relativistic bunch partially cancels the electrostatic repulsion by a factor of $1 - \beta^2 = \gamma^{-2}$, so that the incoherent Laslett space-charge betatron tune shift scales inversely with $\beta\gamma^2$:

$$\Delta Q_{\mathrm{sc}} = -\frac{r_p N_b F_g}{2\pi\epsilon_n \beta\gamma^2 B_f}, \tag{1}$$

where $r_p = q^2/(4\pi\epsilon_0 m_p c^2)$ is the classical proton radius, N_b is the bunch intensity, ϵ_n is the normalized transverse RMS emittance, β and γ are the relativistic Lorentz factors, B_f is the bunching factor, and F_g is the transverse form factor. Consequently, while space-charge tune depression is negligible ($\Delta Q_{\mathrm{sc}} \sim 10^{-8}$–$10^{-9}$) at multi-TeV collider top energies [3], it constitutes the primary intensity bottleneck in low-to-medium energy accumulator rings and injector boosters ($\gamma \sim 1.15$–3.1). Specifically, in machines such as the Los Alamos Proton Storage Ring (PSR, $E_k = 800\,\mathrm{MeV}$, $\gamma = 1.853$, circumference $C = 90.2\,\mathrm{m}$, revolution period $T_{\mathrm{rev}} = 0.358\,\mu\mathrm{s}$), the Spallation Neutron Source (SNS) accumulator ring ($E_k = 1.0\,\mathrm{GeV}$, $\gamma = 2.066$, $C = 248.0\,\mathrm{m}$, $T_{\mathrm{rev}} = 0.945\,\mu\mathrm{s}$) [4], the CERN Proton Synchrotron Booster (PSB, $E_k = 160\,\mathrm{MeV}$–$2.0\,\mathrm{GeV}$, $T_{\mathrm{rev}} = 0.57$–$1.01\,\mu\mathrm{s}$) [5], and the FAIR SIS18 synchrotron [6], the incoherent tune depression reaches $|\Delta Q_{\mathrm{sc}}| \approx 0.1$–$0.4$.

When this large tune shift couples with injection envelope mismatch or pushes particles across low-order structure resonances ($m_x Q_x + m_y Q_y = k$), parametric collective resonances drive particles out of the high-density bunch core into an extended high-action halo [7], causing beam loss and radioactivation of the ring aperture.

Historically, active feedback in circular accelerators has been restricted to dipolar (center-of-mass) centroid damping. Simon van der Meer's stochastic cooling [8, 9], awarded the 1984 Nobel Prize in Physics, damps incoherent phase-space fluctuations over 10^4–10^7 turns. By contrast, coherent envelope oscillations and space-charge-induced mismatch in compact high-intensity rings evolve over a few revolution turns ($T_{\mathrm{rev}} \approx 0.36$–$1.0\,\mu\mathrm{s}$, corresponding to a 5-turn control window of $\leq 5\,\mu\mathrm{s}$). Conventional self-consistent space-charge solvers cannot execute within this microsecond compute budget: direct pairwise Coulomb summation over N mesh or macro-particle nodes scales as $O(N^2)$ ($\sim 1.6\,\mathrm{ms}$ even for $N = 10^3$), while grid-based Fast Fourier Transform (FFT) or multigrid Particle-in-Cell (PIC) codes [10, 11] scale as $O(N_p + G\log G)$ for N_p macro-particles and G grid cells, requiring milliseconds to seconds per step.

In this work, we present the mathematical formalization and open-source numerical verification of a closed-form reduced-order Moment Approach to 3D space-charge field estimation and fast feedback control. It is important to distinguish the two components of this contribution:

1. **1994 LANL Reduced-Order Moment Formulation (Sections 2–4.1)**: Originally motivated by Channell's moment dynamics [12] and initiated during 1994 Los Alamos National Laboratory (LANL) studies for the BEDLAM code [13], the author derived the normalized-prior affine charge inversion ('report1.pdf'), the weighted polynomial potential reduction ('poly-phi.pdf'), and the direct 3D vector electric-field reduction ('poly-e.pdf'),

showing that the mesh dimension N is algebraically eliminated offline into fixed 35×35 and 34×35 matrices.

2. **2026 Moment-Space Zero-Self-Force (ZSF) Dipolar Offset Projection and De-Confounded Closed-Loop Feasibility Study (Sections 4.2–5)**: Extending the 1994 operator to modern real-time feedback, we resolve the non-zero self-force residual that arises when applying a fixed prior-weighted operator ($\boldsymbol{w} = \boldsymbol{\mu}$) to a perturbed beam distribution ($\boldsymbol{a} \neq \boldsymbol{\mu}$) via an exact $O(M_E)$ moment-space dipolar offset projection, and evaluate both coordinate-frame alignment and ZSF projection independently in a reproducible multi-turn 2D-in-3D pairwise-Coulomb Leapfrog tracking and ten-mode ablation benchmark.

—

2 Underdetermined Charge Inversion and Prior Regularization

2.1 Moment Representation in 3D Real Space

In the co-moving rest frame of the bunch, let the spatial domain $\mathcal{V} \subset \mathbb{R}^3$ be discretized into N mesh points $\boldsymbol{r}_i = (x_i, y_i, z_i)$ $(i = 1, \ldots, N)$. Let a_i denote the normalized point-charge weight at mesh node $\boldsymbol{r}_i$, with $\boldsymbol{a} = [a_1, \ldots, a_N]^T \in \mathbb{R}^N$.

The macroscopic geometry of the bunch is characterized by M spatial moments up to polynomial order $P = 4$:

$$s_k = \langle m_k(\boldsymbol{r}) \rangle_{\boldsymbol{a}} = \sum_{i=1}^{N} m_k(\boldsymbol{r}_i) a_i, \quad k = 1, \ldots, M, \tag{2}$$

where $m_k(\boldsymbol{r}) = x^{p_k} y^{q_k} z^{l_k}$ $(p_k + q_k + l_k \leq 4)$ are the $M = \binom{4+3}{3} = 35$ distinct 3D monomial basis functions:

$$\boldsymbol{m}(\boldsymbol{r}) = \left[1, x, y, z, x^2, xy, xz, y^2, yz, z^2, x^3, \ldots, z^4\right]^T. \tag{3}$$

Defining the moment evaluation matrix $\boldsymbol{\Omega} \in \mathbb{R}^{M \times N}$ with entries $\Omega_{ki} = m_k(\boldsymbol{r}_i)$, Eq. (2) takes the linear form:

$$\boldsymbol{s} = \boldsymbol{\Omega} \boldsymbol{a}. \tag{4}$$

When $N \gg M$ (typically $N \sim 10^3$–10^5 while $M = 35$) and the mesh points are non-coplanar on algebraic surfaces of degree ≤ 4 such that $\operatorname{rank}(\boldsymbol{\Omega}) = M = 35$, Eq. (4) is an underdetermined linear inverse problem with an $(N - M)$-dimensional nullspace.

2.2 Constrained Variational Inversion and Fringe Negativity Suppression

To select a unique, smooth mesh charge assignment $\boldsymbol{a}$ without iterative non-negative quadratic programming (which would violate microsecond real-time constraints), we minimize the constrained quadratic functional [13]:

$$\mathcal{J}(\boldsymbol{a}, \boldsymbol{\lambda}) = \frac{1}{2} \sum_{i=1}^{N} (a_i - \mu_i)^2 + \boldsymbol{\lambda}^T (\boldsymbol{s} - \boldsymbol{\Omega} \boldsymbol{a}), \tag{5}$$

where $\boldsymbol{\mu} = [\mu_1, \ldots, \mu_N]^T \in \mathbb{R}^N$ is a fixed reference prior distribution and $\boldsymbol{\lambda} \in \mathbb{R}^M$ is the Lagrange multiplier vector.

In early tests for the BEDLAM code at LANL, setting $\boldsymbol{\mu} = \boldsymbol{0}$ yielded the minimum-L_2-norm pseudoinverse solution $\boldsymbol{a}_{\min} = \boldsymbol{\Omega}^T (\boldsymbol{\Omega} \boldsymbol{\Omega}^T)^{-1} \boldsymbol{s}$, which exhibits large unphysical negative charge oscillations ($a_i < 0$) in the low-density bunch fringe whenever the bunch radius is smaller than the mesh radius ($R_{\text{beam}} < R_{\text{mesh}}$). Paul Channell suggested a uniform reference prior $\mu_i = \text{const}$ [12].

However, because the zeroth-order constant monomial $m_1(\boldsymbol{r}) = 1$ lies in the row space of $\boldsymbol{\Omega}$, the orthogonal projection operator $[\boldsymbol{I}_N - \boldsymbol{\Omega}^T(\boldsymbol{\Omega}\boldsymbol{\Omega}^T)^{-1}\boldsymbol{\Omega}]$ annihilates any constant vector identically ($\|\boldsymbol{a}(\mu = \text{const}) - \boldsymbol{a}(\mu = 0)\|_\infty = 0$), leaving the fringe negativity unchanged.

As shown by the author in the April 1994 LANL study [13], effective non-iterative regularization requires choosing a spatially localized physical reference profile $\boldsymbol{\mu}$ whose second moments match the nominal beam envelope and whose total weight is normalized to unity:

$$\sum_{i=1}^{N} \mu_i = 1, \quad \mu_i > 0 \quad \forall i. \tag{6}$$

Theorem 1 (Closed-Form Affine Charge Reconstruction)**.** *Assume the N-point 3D mesh satisfies the full-row-rank condition* $\text{rank}(\boldsymbol{\Omega}) = M = 35$*, so that $\boldsymbol{\Omega}\boldsymbol{\Omega}^T \in \mathbb{R}^{35\times 35}$ is strictly symmetric positive-definite. Under the normalization constraint (6), the unique global minimizer of $\mathcal{J}(\boldsymbol{a}, \boldsymbol{\lambda})$ subject to $\boldsymbol{\Omega}\boldsymbol{a} = \boldsymbol{s}$ is the affine transformation:*

$$\boldsymbol{a} = \boldsymbol{\mu} + \boldsymbol{\Lambda}(\boldsymbol{s} - \boldsymbol{\Omega}\boldsymbol{\mu}) = \boldsymbol{\Lambda}\boldsymbol{s} + \boldsymbol{c}, \tag{7}$$

where $\boldsymbol{\Lambda} \in \mathbb{R}^{N\times M}$ is the Moore-Penrose right pseudoinverse and $\boldsymbol{c} \in \mathbb{R}^N$ is the prior residual vector:

$$\boldsymbol{\Lambda} = \boldsymbol{\Omega}^T(\boldsymbol{\Omega}\boldsymbol{\Omega}^T)^{-1}, \qquad \boldsymbol{c} = (\boldsymbol{I}_N - \boldsymbol{\Lambda}\boldsymbol{\Omega})\boldsymbol{\mu}. \tag{8}$$

At the reference equilibrium ($\boldsymbol{s} = \boldsymbol{\Omega}\boldsymbol{\mu}$), Eq. (7) yields $\boldsymbol{a} \equiv \boldsymbol{\mu} > \boldsymbol{0}$, guaranteeing strict pointwise positivity across all N mesh points. For perturbed beams ($\boldsymbol{s} \neq \boldsymbol{\Omega}\boldsymbol{\mu}$), Eq. (7) preserves all $M = 35$ moments ($\boldsymbol{\Omega}\boldsymbol{a} = \boldsymbol{s}$) identically, and maintains pointwise positivity ($a_i > 0$) if and only if the moment perturbation satisfies the componentwise bound $[\boldsymbol{\Lambda}(\boldsymbol{s} - \boldsymbol{\Omega}\boldsymbol{\mu})]_i > -\mu_i$.

Proof. Setting $\nabla_{\boldsymbol{a}}\mathcal{J} = \boldsymbol{0}$ gives $\boldsymbol{a} - \boldsymbol{\mu} - \boldsymbol{\Omega}^T\boldsymbol{\lambda} = \boldsymbol{0}$, or $\boldsymbol{a} = \boldsymbol{\mu} + \boldsymbol{\Omega}^T\boldsymbol{\lambda}$. Multiplying from the left by $\boldsymbol{\Omega}$ and enforcing $\boldsymbol{\Omega}\boldsymbol{a} = \boldsymbol{s}$ yields $\boldsymbol{\Omega}\boldsymbol{\mu} + (\boldsymbol{\Omega}\boldsymbol{\Omega}^T)\boldsymbol{\lambda} = \boldsymbol{s}$. Because $\text{rank}(\boldsymbol{\Omega}) = M$, $\boldsymbol{\Omega}\boldsymbol{\Omega}^T$ is invertible, giving $\boldsymbol{\lambda} = (\boldsymbol{\Omega}\boldsymbol{\Omega}^T)^{-1}(\boldsymbol{s} - \boldsymbol{\Omega}\boldsymbol{\mu})$. Substituting $\boldsymbol{\lambda}$ back into $\boldsymbol{a}$ establishes Eq. (7), and the positivity condition $a_i > 0 \iff [\boldsymbol{\Lambda}(\boldsymbol{s} - \boldsymbol{\Omega}\boldsymbol{\mu})]_i > -\mu_i$ follows immediately. □

Remark on Finite Perturbations Outside the Positivity Neighborhood: Outside the perturbative neighborhood $[\boldsymbol{\Lambda}(\boldsymbol{s} - \boldsymbol{\Omega}\boldsymbol{\mu})]_i > -\mu_i$, unconstrained affine inversion permits localized negative weights ($a_i < 0$) in the low-density fringe when $R_{\text{beam}} < R_{\text{mesh}}$. However, because $\boldsymbol{c} = (\boldsymbol{I}_N - \boldsymbol{\Lambda}\boldsymbol{\Omega})\boldsymbol{\mu}$ supplies the positive nullspace component of the localized reference profile $\boldsymbol{\mu}$, empirical evaluation on our 3D anisotropic bunch benchmark (Section 5.2) demonstrates that Eq. (7) reduces integrated fringe negative-charge overshoot by 3.04× (0.2669 → 0.0876) and L_2 reconstruction error by 2.00× relative to the uniform/zero prior.

—

3 Space-Charge Potential Reduction and the Fall-Out of N

3.1 Interior Polynomial Potential Fitting via Weighted Least Squares

Given a mesh charge distribution $\boldsymbol{a}$ in the bunch rest frame, the regularized discrete electrostatic potential on the mesh is given by the Plummer-softened Coulomb sum:

$$\phi_{\text{exact}}(\boldsymbol{r}_i) = \sum_{j=1}^{N} D_{ij}a_j, \qquad D_{ij} = \begin{cases} \dfrac{1}{\sqrt{|\boldsymbol{r}_i - \boldsymbol{r}_j|^2 + h^2}}, & i \neq j, \\ 0, & i = j, \end{cases} \tag{9}$$

where $h = \alpha_h N^{-1/3}$ is the inter-particle mesh softening scale that prevents singular two-body collisions on the discrete grid and $D_{ii} = 0$ excludes self-interaction.

To represent the interior space-charge potential analytically within the beam core, we expand ϕ in the 4th-order monomial basis:

$$\phi_{\text{approx}}(\boldsymbol{r}) = \sum_{l=1}^{M} v_l m_l(\boldsymbol{r}) = \boldsymbol{m}^T(\boldsymbol{r})\boldsymbol{v}, \tag{10}$$

where $\boldsymbol{v} = [v_1, \ldots, v_M]^T \in \mathbb{R}^M$. Because a 4th-order polynomial grows as r^4 at large radii whereas the exterior Laplace potential decays as r^{-1}, an unweighted fit over the bounding box would be distorted by empty vacuum nodes outside the bunch. We therefore weight the least-squares functional by a positive density measure w_i concentrated in the beam core:

$$\chi^2_{\boldsymbol{w}}(\boldsymbol{v}) = \sum_{i=1}^{N} w_i \left[\phi_{\text{exact}}(\boldsymbol{r}_i) - \phi_{\text{approx}}(\boldsymbol{r}_i)\right]^2. \tag{11}$$

Theorem 2 (Potential Dimension Reduction and Weighted Moment Conservation)**.** *For any weight vector $\boldsymbol{w} \in \mathbb{R}^N$ such that the weighted Gram matrix $\boldsymbol{\Omega}\{\boldsymbol{\Omega w}\}^T \in \mathbb{R}^{M\times M}$ is non-singular (which holds whenever $w_i > 0$ on a rank-M mesh), the normal equations minimizing $\chi^2_{\boldsymbol{w}}(\boldsymbol{v})$ yield the unique coefficient vector:*

$$\boldsymbol{v}(\boldsymbol{w}) = \boldsymbol{\Lambda}'^T(\boldsymbol{w})\boldsymbol{\phi}_{\text{exact}}, \qquad \boldsymbol{\Lambda}'(\boldsymbol{w}) = \{\boldsymbol{\Omega w}\}^T\left[\boldsymbol{\Omega}\{\boldsymbol{\Omega w}\}^T\right]^{-1}, \tag{12}$$

where $\{\boldsymbol{\Omega w}\}_{ki} = \Omega_{ki} w_i$. By virtue of the normal equations, the polynomial potential $\boldsymbol{\phi}_{\text{approx}} = \boldsymbol{\Omega}^T\boldsymbol{v}(\boldsymbol{w})$ preserves all $M = 35$ $\boldsymbol{w}$-weighted potential moments identically:

$$\langle m_k \phi_{\text{approx}}\rangle_{\boldsymbol{w}} = \sum_{i=1}^{N} w_i m_k(\boldsymbol{r}_i)\phi_{\text{approx}}(\boldsymbol{r}_i) = \langle m_k \phi_{\text{exact}}\rangle_{\boldsymbol{w}}, \qquad k = 1, \ldots, M. \tag{13}$$

Furthermore, adopting the fixed reference-prior weighting $\boldsymbol{w} = \boldsymbol{\mu}$ (first-order perturbative linearization around $\boldsymbol{a} = \boldsymbol{\mu}$) eliminates the mesh dimension N entirely offline, yielding an algebraically exact reduced-order representation of the $\boldsymbol{\mu}$-weighted polynomial projection on the reconstructed charge $\boldsymbol{a} = \boldsymbol{\Lambda s} + \boldsymbol{c}$:

$$\boldsymbol{v} = \boldsymbol{K s} + \boldsymbol{v}_0, \qquad \boldsymbol{K} = \boldsymbol{\Lambda}'^T(\boldsymbol{\mu})\boldsymbol{D}\boldsymbol{\Lambda} \in \mathbb{R}^{M\times M}, \qquad \boldsymbol{v}_0 = \boldsymbol{\Lambda}'^T(\boldsymbol{\mu})\boldsymbol{D}\boldsymbol{c} \in \mathbb{R}^M. \tag{14}$$

Proof. Setting $\partial\chi^2_{\boldsymbol{w}}/\partial v_k = 0$ for $k = 1, \ldots, M$ yields the weighted normal equations:

$$\sum_{i=1}^{N} w_i \Omega_{ki} \left(\boldsymbol{\Omega}^T\boldsymbol{v}\right)_i = \sum_{i=1}^{N} w_i \Omega_{ki}\phi_{\text{exact}}(\boldsymbol{r}_i) \iff \{\boldsymbol{\Omega w}\}\boldsymbol{\Omega}^T\boldsymbol{v} = \{\boldsymbol{\Omega w}\}\boldsymbol{\phi}_{\text{exact}}, \tag{15}$$

which establishes the moment conservation identity (13) and, upon left-multiplying by $[\{\boldsymbol{\Omega w}\}\boldsymbol{\Omega}^T]^{-1}$, gives $\boldsymbol{v}(\boldsymbol{w}) = \boldsymbol{\Lambda}'^T(\boldsymbol{w})\boldsymbol{\phi}_{\text{exact}}$. Substituting $\boldsymbol{\phi}_{\text{exact}} = \boldsymbol{D a} = \boldsymbol{D}(\boldsymbol{\Lambda s} + \boldsymbol{c})$ with fixed weight $\boldsymbol{w} = \boldsymbol{\mu}$ factors all N-dimensional matrices into the offline constants $\boldsymbol{K}$ and $\boldsymbol{v}_0$ of Eq. (14). □

Remark on Algebraic Reduction Exactness vs. Physical Field Approximation: It is essential to distinguish two levels of accuracy in Theorem 2: (i) *algebraic exactness of the reduced operator*, meaning that $\boldsymbol{v} = \boldsymbol{K s} + \boldsymbol{v}_0 \in \mathbb{R}^{35}$ reproduces the full N-mesh prior-weighted projection $\boldsymbol{\Lambda}'^T(\boldsymbol{\mu})\boldsymbol{D}(\boldsymbol{\Lambda s} + \boldsymbol{c})$ with zero truncation error (1.14×10^{-13} numerically), and (ii) *physical field approximation*, which involves replacing the true continuum density by the affine reconstruction $\boldsymbol{a}$, truncating the potential basis at degree $P = 4$, and replacing the state-dependent weight $\boldsymbol{w} = \boldsymbol{a}$ by the fixed prior weight $\boldsymbol{w} = \boldsymbol{\mu}$ (introducing an $O(\|\boldsymbol{a} - \boldsymbol{\mu}\|)$ difference relative to the state-weighted fit).

—

4 Direct 3D Vector Electric Field Reduction and Moment-Space Zero-Self-Force Dipolar Offset Projection

4.1 Direct Vector Field Least-Squares Formulation

In beam dynamics, betatron tune shifts and multipole kicks depend directly on the electric field $\boldsymbol{E} = -\nabla\phi$. Rather than differentiating a fitted scalar potential (which amplifies high-order boundary oscillations), the author derived a direct 3D vector electric-field least-squares fit in 1994 ('poly-e.pdf').

Using the Plummer-softened field kernels $G_{x_j,ik} = (x_{j,i} - x_{j,k})(|\boldsymbol{r}_i - \boldsymbol{r}_k|^2 + h^2)^{-3/2}$ ($i \neq k$, $G_{x_j,ii} = 0$), the exact discrete electric field components along x_1, x_2, x_3 (x, y, z) are:

$$\boldsymbol{E}_{x_j}^{\text{exact}} = \boldsymbol{G}_{x_j}\boldsymbol{a}, \qquad j = 1, 2, 3. \tag{16}$$

Because the constant monomial $m_1(\boldsymbol{r}) = 1$ has zero gradient ($\nabla m_1 \equiv \boldsymbol{0}$), we retain only the $M_E = M - 1 = 34$ non-constant monomials ($1 \leq p_l + q_l + l_l \leq 4$), denoted by the reduced derivative matrices $\tilde{\boldsymbol{\Omega}}^{x_j} \in \mathbb{R}^{34\times N}$ with entries $\tilde{\Omega}_{li}^{x_j} = \partial m_l(\boldsymbol{r}_i)/\partial x_j$. The polynomial electric field approximation is:

$$\boldsymbol{E}_{x_j}^{\text{approx}} = -\tilde{\boldsymbol{\Omega}}^{x_j T}\boldsymbol{v}_E, \qquad \boldsymbol{v}_E \in \mathbb{R}^{34}. \tag{17}$$

Minimizing the density-weighted 3D vector residual functional:

$$\mathcal{S}_{\boldsymbol{w}}(\boldsymbol{v}_E) = \sum_{i=1}^{N} w_i \sum_{j=1}^{3} \left[E_{x_j}^{\text{exact}}(\boldsymbol{r}_i) - E_{x_j}^{\text{approx}}(\boldsymbol{r}_i)\right]^2, \tag{18}$$

with respect to $\boldsymbol{v}_E$ yields the following theorem.

Theorem 3 (Direct 3D Electric Field Reduction and Weighted Force Identity)**.** *For any weight vector $\boldsymbol{w} \in \mathbb{R}^N$ such that $\sum_{j=1}^{3}\{\tilde{\boldsymbol{\Omega}}^{x_j}\boldsymbol{w}\}\tilde{\boldsymbol{\Omega}}^{x_j T} \in \mathbb{R}^{34\times 34}$ is invertible, the optimal 34-component electric-field coefficient vector minimizing $\mathcal{S}_{\boldsymbol{w}}(\boldsymbol{v}_E)$ is:*

$$\boldsymbol{v}_E(\boldsymbol{w}) = -\left[\sum_{j=1}^{3}\{\tilde{\boldsymbol{\Omega}}^{x_j}\boldsymbol{w}\}\tilde{\boldsymbol{\Omega}}^{x_j T}\right]^{-1}\left[\sum_{j=1}^{3}\{\tilde{\boldsymbol{\Omega}}^{x_j}\boldsymbol{w}\}\boldsymbol{E}_{x_j}^{\text{exact}}\right]. \tag{19}$$

This solution satisfies all $M_E = 34$ $\boldsymbol{w}$-weighted generalized field-moment normal equations identically:

$$\sum_{j=1}^{3}\left\langle \frac{\partial m_k}{\partial x_j} E_{x_j}^{\text{approx}}(\boldsymbol{w})\right\rangle_{\boldsymbol{w}} = \sum_{j=1}^{3}\left\langle \frac{\partial m_k}{\partial x_j} E_{x_j}^{\text{exact}}\right\rangle_{\boldsymbol{w}}, \qquad k = 2, \ldots, 35. \tag{20}$$

*In particular, evaluating Eq. (20) for the three linear dipolar monomials $m_k \in \{x, y, z\}$ (where $\partial x_i/\partial x_j = \delta_{ij}$) guarantees the $\boldsymbol{w}$-weighted **Average Force Identity**:*

$$\langle \boldsymbol{E}_{\text{approx}}(\boldsymbol{w})\rangle_{\boldsymbol{w}} = \langle \boldsymbol{E}_{\text{exact}}\rangle_{\boldsymbol{w}}. \tag{21}$$

Consequently:

1. ***State-Weighted Fit ($\boldsymbol{w} = \boldsymbol{a}$)**: By the antisymmetry $G_{x_j,ik} = -G_{x_j,ki}$ (Newton's third law), $\langle E_{x_j}^{\text{exact}}\rangle_{\boldsymbol{a}} = \boldsymbol{a}^T\boldsymbol{G}_{x_j}\boldsymbol{a} = 0$, so the state-weighted fit satisfies $\langle \boldsymbol{E}_{\text{approx}}(\boldsymbol{a})\rangle_{\boldsymbol{a}} = \boldsymbol{0}$ identically.*
2. ***Prior-Weighted Offline Reduction ($\boldsymbol{w} = \boldsymbol{\mu}$)**: Fixing $\boldsymbol{w} = \boldsymbol{\mu}$ collapses Eq. (19) offline into the runtime 34×35 affine operator:*

$$\boldsymbol{v}_E^{\text{prior}} = \boldsymbol{K}_E\boldsymbol{s} + \boldsymbol{v}_{E,0}, \qquad \boldsymbol{K}_E = \boldsymbol{\Lambda}_E^T(\boldsymbol{\mu})\boldsymbol{\Lambda} \in \mathbb{R}^{34\times 35}, \qquad \boldsymbol{v}_{E,0} = \boldsymbol{\Lambda}_E^T(\boldsymbol{\mu})\boldsymbol{c} \in \mathbb{R}^{34}, \tag{22}$$

where $\boldsymbol{\Lambda}_E^T(\boldsymbol{\mu}) = -\left[\sum_{j=1}^{3}\{\tilde{\boldsymbol{\Omega}}^{x_j}\boldsymbol{\mu}\}\tilde{\boldsymbol{\Omega}}^{x_j T}\right]^{-1}\left(\sum_{j=1}^{3}\{\tilde{\boldsymbol{\Omega}}^{x_j}\boldsymbol{\mu}\}\boldsymbol{G}_{x_j}\right) \in \mathbb{R}^{34\times N}$. For a perturbed state $\boldsymbol{a} = \boldsymbol{\mu} + \delta\boldsymbol{a}$, Eq. (22) preserves the $\boldsymbol{\mu}$-weighted field averages $\langle E_{x_j}^{\text{approx}}(\boldsymbol{\mu})\rangle_{\boldsymbol{\mu}} = \langle E_{x_j}^{\text{exact}}\rangle_{\boldsymbol{\mu}} = \boldsymbol{\mu}^T\boldsymbol{G}_{x_j}\delta\boldsymbol{a} \neq 0$, which generally leaves a non-zero residual net force $\langle \boldsymbol{E}_{\text{approx}}(\boldsymbol{\mu})\rangle_{\boldsymbol{a}} \neq \boldsymbol{0}$ on the actual distribution $\boldsymbol{a}$.

4.2 Exact Moment-Space Zero-Self-Force (ZSF) Dipolar Offset Projection

As highlighted in Theorem 3, because $\boldsymbol{w} = \boldsymbol{\mu}$ is fixed offline, the uncorrected prior-weighted operator $\boldsymbol{v}_E^{\text{prior}}$ does not automatically vanish when averaged over a perturbed charge state $\boldsymbol{a} \neq \boldsymbol{\mu}$: both $\boldsymbol{\mu}^T \boldsymbol{G}_{x_j} \delta\boldsymbol{a} \neq 0$ and the weighting mismatch $(\boldsymbol{a} - \boldsymbol{\mu})^T \boldsymbol{E}_{x_j}^{\text{approx}}(\boldsymbol{\mu}) \neq 0$ contribute a spurious dipolar self-force $\langle \boldsymbol{E}_{\text{approx}}(\boldsymbol{\mu}) \rangle_{\boldsymbol{a}} \neq \boldsymbol{0}$.

Crucially, because the spatial derivative of every monomial of degree ≤ 4 is a monomial of degree ≤ 3 already contained in $\boldsymbol{m}(\boldsymbol{r})$, there exist exact constant differentiation matrices $\boldsymbol{D}_{x_j}^{\text{mom}} \in \mathbb{R}^{34\times 35}$ such that $\tilde{\boldsymbol{\Omega}}^{x_j} \equiv \boldsymbol{D}_{x_j}^{\text{mom}} \boldsymbol{\Omega}$. Therefore, for any polynomial coefficient vector $\boldsymbol{v}_E \in \mathbb{R}^{34}$ and any charge state $\boldsymbol{a}$ with moments $\boldsymbol{s} = \boldsymbol{\Omega}\boldsymbol{a}$, the instantaneous state-weighted net self-force depends *exclusively on the* 35*-element moment vector* $\boldsymbol{s}$:

$$\langle E_{x_j}^{\text{approx}} \rangle_{\boldsymbol{a}} = \boldsymbol{a}^T \left(-\tilde{\boldsymbol{\Omega}}^{x_j T} \boldsymbol{v}_E \right) = -\left(\boldsymbol{D}_{x_j}^{\text{mom}} \boldsymbol{\Omega}\boldsymbol{a} \right)^T \boldsymbol{v}_E = -\left(\boldsymbol{D}_{x_j}^{\text{mom}} \boldsymbol{s} \right)^T \boldsymbol{v}_E. \tag{23}$$

Corollary 4 (Moment-Space Zero-Self-Force (ZSF) Dipolar Offset Projection)**.** *Let* $\boldsymbol{v}_E^{\text{prior}} = \boldsymbol{K}_E \boldsymbol{s} + \boldsymbol{v}_{E,0} \in \mathbb{R}^{34}$ *be the prior-weighted field coefficients, and let* $l_x, l_y, l_z \in \{1, \ldots, 34\}$ *denote the indices of the three linear dipolar basis monomials* (x, y, z)*, for which* $(\boldsymbol{D}_{x_j}^{\text{mom}} \boldsymbol{s})_{l_{x_k}} = \delta_{jk} s_{000}$*. Defining the ZSF-projected coefficient vector* $\boldsymbol{v}_E^{\text{ZSF}} \in \mathbb{R}^{34}$ *by leaving all* 31 *higher-order multipole coefficients* $(p + q + l \geq 2)$ *unchanged and setting the three spatially uniform dipolar components to:*

$$v_{E,l_{x_j}}^{\text{ZSF}} = -\frac{1}{s_{000}} \sum_{\substack{l=1 \\ l \neq l_{x_j}}}^{34} v_{E,l}^{\text{prior}} \left(\boldsymbol{D}_{x_j}^{\text{mom}} \boldsymbol{s} \right)_l, \qquad j = 1, 2, 3, \tag{24}$$

enforces exact zero net self-force on the instantaneous perturbed distribution $\boldsymbol{a}$*:*

$$\langle \boldsymbol{E}_{\text{approx}}^{\text{ZSF}} \rangle_{\boldsymbol{a}} = \langle \boldsymbol{E}_{\text{exact}} \rangle_{\boldsymbol{a}} = \boldsymbol{0}, \tag{25}$$

in $O(M_E)$ *scalar operations (*57 *non-zero derivative terms across* x, y, z*) using only* $\boldsymbol{s} \in \mathbb{R}^{35}$*, without accessing the* N*-point mesh or altering any quadrupolar, sextupolar, or octupolar field gradient. (We refer to Eq.* (24) *as a* dipolar offset projection *or* zero-mean-force projection*, since adding a linear term* $\delta\boldsymbol{c} \cdot \boldsymbol{r}$ *to* Φ_{poly} *shifts the spatially uniform dipolar field component* $-\delta\boldsymbol{c}$ *to satisfy Newton's third law, rather than shifting a constant scalar potential gauge.)*

—

5 Hardware Architecture and Numerical Verification

5.1 Moment Acquisition, Actuator Realizability, and Latency Budget

In a physical synchrotron or accumulator ring, applying Eq. (22) and Corollary 4 within a closed-loop feedback architecture involves three practical system constraints:

1. **Hybrid Diagnostic-Observer Pipeline**: Standard wall-current monitors and 4-button/stripline Beam Position Monitors (BPMs) directly measure bunch intensity s_{000}, centroid offsets ($\langle x \rangle, \langle y \rangle, \langle z \rangle$), and the transverse quadrupolar difference moment $\langle x^2 - y^2 \rangle$ (plus the skew-quadrupolar moment $\langle xy \rangle$ when equipped with rotated or 8-electrode pickups), whereas the quadrupolar sum $\langle x^2 + y^2 \rangle$ and 3rd- and 4th-order moments $(3 \leq p + q + l \leq 4)$ are not directly observable from a single standard button BPM and require a synchronized online envelope/moment observer constrained by multi-BPM phase-advance tomography or periodic profile measurements.

2. **Magnetic Corrector Realizability ($\boldsymbol{c}_{\text{achievable}} = \boldsymbol{B}\boldsymbol{u}$)**: The $M_E = 34$ polynomial coefficients $\boldsymbol{v}_E^{\text{ZSF}}$ parameterize an interior 3D electrostatic field ($\nabla \cdot \boldsymbol{E} \neq 0$ inside the bunch charge), whereas external magnetic correctors apply localized 2D transverse harmonic kicks satisfying Maxwell's vacuum equations ($\nabla \cdot \boldsymbol{B}_{\text{ext}} = 0, \nabla \times \boldsymbol{B}_{\text{ext}} = \boldsymbol{0}$). Thus, a physical corrector family with control currents $\boldsymbol{u}$ spans an achievable subspace $\boldsymbol{c}_{\text{achievable}} = \boldsymbol{B}\boldsymbol{u}$. In our ablation study below, we explicitly contrast **Modes B-Lab, B-CF, and B-CF-ZSF**—which restrict actuation to the 11 low-order dipole, quadrupole, and octupole terms corresponding to standard corrector magnet families—against **Modes C-Lab, C-CF, and C-CF-ZSF**, which apply all 34 polynomial basis terms as a theoretical upper bound on 4th-order moment compensation.

3. **Compute Kernel Latency vs. Total Closed-Loop Latency**: Total hardware loop delay is the sum of diagnostic acquisition, observer estimation, matrix computation, communication, and magnet power-supply/vacuum-chamber eddy-current response:

$$T_{\text{total}} = T_{\text{diag}} + T_{\text{obs}} + T_{\text{compute}} + T_{\text{comm}} + T_{\text{act}}. \tag{26}$$

On a single CPU core, the 35×35 and 34×35 matrix-vector kernels execute in $\sim$ 1.6–5.8 μs in NumPy/OpenBLAS (T_{compute} only); adding Python-level ZSF array indexing requires $\sim$ 5.4–20 μs (hardware- and load-dependent; the companion benchmark (`benchmark_accelerator_space_charge.py`) reports values in this range), consistently exceeding the 4.73 μs 5-turn window of the SNS ring in interpreted Python. All quoted latencies refer exclusively to T_{compute} of Eq. (26); end-to-end feedback latency T_{total} additionally includes T_{diag}, T_{obs}, T_{comm}, and T_{act}, which were not demonstrated in this work. Because the entire runtime operator consists of $34 \times 35 = 1,190$ matrix-vector multiply-accumulate (MAC) operations plus 57 sparse ZSF derivative MACs (1,247 MACs total), a pipelined C or FPGA implementation is projected to execute T_{compute} in $< 1.0\,\mu$s (subject to hardware synthesis, timing closure, and memory access characterization), suggesting that the reduced-order space-charge field solver would not be the computational bottleneck of a prospective feedback loop.

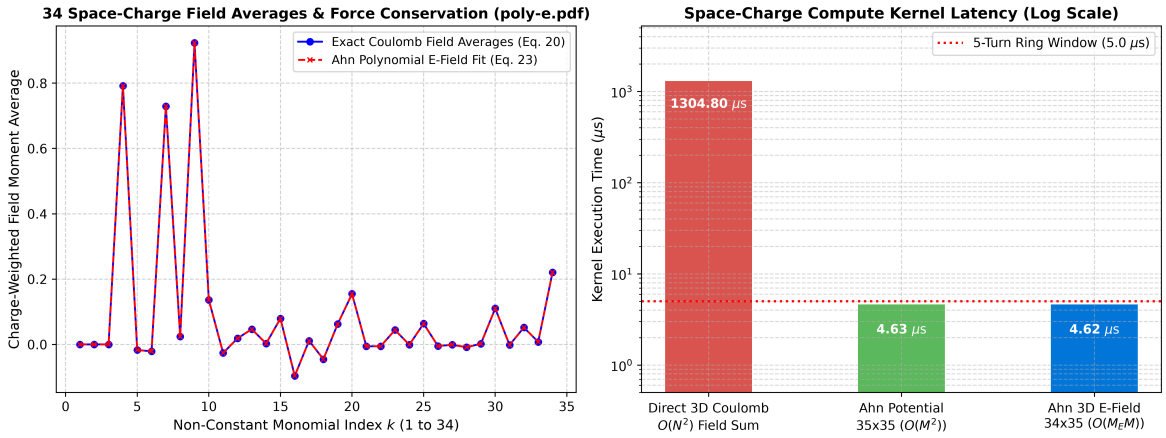


Figure 1: Numerical verification and compute-kernel latency benchmark of the High-Intensity Accelerator Space-Charge Engine (`accelerator_space_charge_engine.py`). *Left*: Exact agreement across all $M_E = 34$ generalized electric-field moment averages between the exact 3D Coulomb field (Eq. (16)) and the state-weighted polynomial field fit (Eq. (19)), confirming Theorem 3 ($\|\langle \boldsymbol{E}_{\text{approx}} \rangle_{\boldsymbol{a}} - \langle \boldsymbol{E}_{\text{exact}} \rangle_{\boldsymbol{a}}\|_\infty = 4.58 \times 10^{-16}$). *Right*: Log-scale compute-kernel latency (T_{compute}) comparing direct $O(N^2)$ 3D Coulomb summation (1595 μs) against the closed-form 35×35 potential kernel and 34×35 3D electric-field kernel.

5.2 Verification of Theorems 1–3, ZSF Dipolar Projection, and Core Fidelity

We verify all mathematical claims using the companion open-source benchmark suite (`benchmark_accelerator_space_charge.py`) on a 3D spherical mesh of $N = 1,000$ points ($R_{\text{mesh}} = 1.0$, Plummer softening $h = 0.05$, full basis rank $\text{rank}(\boldsymbol{\Omega}) = 35/35$) for a 1.0 GeV high-intensity proton accumulator/booster ring ($\gamma = 2.066, \beta = 0.875, C = 248.0\,\text{m}, T_{\text{rev}} = 0.945\,\mu\text{s}$, 5-turn control window $= 4.73\,\mu\text{s}$, Laslett tune depression $\Delta Q_{\text{sc}} = -0.2271$):

- **Theorem 1 (Normalized Prior & Moment Inversion)**: At the reference equilibrium ($\boldsymbol{s} = \boldsymbol{\Omega}\boldsymbol{\mu}$), reconstructed charges exhibit strictly zero negative nodes (0/1000). For an anisotropic, off-axis parabolic bunch ($R_{\text{beam}} \approx 0.55\,R_{\text{mesh}}$), Channell's uniform prior $\mu_i = 1/N$ is numerically identical to the zero prior ($\|\boldsymbol{a}(\mu = \text{const}) - \boldsymbol{a}(\mu = 0)\|_\infty = 8.87 \times 10^{-17}$), whereas our normalized physical prior reduces total fringe negative-charge overshoot by 3.04× (0.2669 → 0.0876) and L_2 profile error by 2.00× (0.0326 → 0.0163), while conserving all 35 moments to $\|\boldsymbol{\Omega}\boldsymbol{a} - \boldsymbol{s}\|_\infty = 1.44 \times 10^{-15}$.

- **Theorem 2 (Potential Reduction)**: The closed-form 35×35 operator $\boldsymbol{v} = \boldsymbol{K}\boldsymbol{s} + \boldsymbol{v}_0$ matches the full N-grid prior-weighted projection $\boldsymbol{\Lambda}'^T(\boldsymbol{\mu})\boldsymbol{D}\boldsymbol{a}$ to 1.14×10^{-13}, and preserves all 35 weighted potential moments to 4.44×10^{-16}.

- **Theorem 3 & Corollary 4 (Comparison of All Three Field Operators)**: Within the high-density bunch core ($r \le 0.45$, containing 79.9% of total positive bunch weight), we evaluate the charge-weighted field fit R^2, relative RMS error ϵ_E, and net self-force $\|\langle \boldsymbol{E}_{\text{approx}} \rangle_{\boldsymbol{a}}\|_\infty$ across all three operators:

  1. **State-Weighted Fit ($\boldsymbol{w} = \boldsymbol{a}$, $O(N^2)$ reference)**: Preserves all 34 $\boldsymbol{a}$-weighted field averages to 4.17×10^{-16}, yields exact zero self-force $\|\langle \boldsymbol{E}_{\text{approx}}(\boldsymbol{a}) \rangle_{\boldsymbol{a}}\|_\infty = 4.58 \times 10^{-16}$, and achieves a core field fit of $R^2 = 0.8962$ ($\epsilon_E = 32.22\%$).
  2. **Uncorrected Prior Runtime Operator ($\boldsymbol{w} = \boldsymbol{\mu}$, $\boldsymbol{v}_E^{\text{prior}} = \boldsymbol{K}_E\boldsymbol{s} + \boldsymbol{v}_{E,0}$)**: Preserves the $\boldsymbol{\mu}$-weighted field averages to 2.03×10^{-15} and matches the N-grid projection $\boldsymbol{\Lambda}_E^T(\boldsymbol{\mu})\boldsymbol{a}$ to 1.07×10^{-14} with $R^2 = 0.8868$ ($\epsilon_E = 33.64\%$) and quadrupole gradient agreement within 11.74%, but leaves a non-zero net force on the perturbed state $\boldsymbol{a}$: $\langle \boldsymbol{E}_{\text{approx}}^{\text{prior}} \rangle_{\boldsymbol{a}} = (0.1123, -0.0732, 0.0727)$ ($\|\langle \boldsymbol{E}_{\text{approx}}^{\text{prior}} \rangle_{\boldsymbol{a}}\|_\infty = 0.1123$).
  3. **ZSF Dipolar-Projected Runtime Operator ($\boldsymbol{w} = \boldsymbol{\mu} + \text{ZSF}$, $\boldsymbol{v}_E^{\text{ZSF}}$)**: Restores exact zero net self-force on the perturbed distribution $\boldsymbol{a}$ to machine precision ($\|\langle \boldsymbol{E}_{\text{approx}}^{\text{ZSF}} \rangle_{\boldsymbol{a}}\|_\infty = 4.02 \times 10^{-16}$, or 9.02×10^{-17} on the true charge vector) while preserving all 31 higher-order multipole gradients and slightly improving the core field fit to $R^2 = 0.8878$ ($\epsilon_E = 33.50\%$).

5.3 Idealized Multi-Turn 2D-in-3D Pairwise-Coulomb Tracking and De-Confounded Ten-Mode Ablation Study

To evaluate closed-loop multipole compensation without circular toy models or initial-condition rescaling, and to cleanly separate the effect of *centroid-frame coordinate alignment* from *Corollary 4's ZSF dipolar offset projection*, we execute an idealized zero-latency 2D transverse (x, p_x, y, p_y) Leapfrog tracking and de-confounded ten-mode ablation study in a 3D pairwise-Coulomb bunch geometry (`simulate_multiturn_envelope_tracking`). Because the 80-turn betatron tracking window is much shorter than a synchrotron oscillation period ($80 \ll T_{\text{synch}} \sim 10^3$ turns), longitudinal particle coordinates z_p are held quasi-static during the betatron window while acting as full 3D sources in the pairwise Coulomb sum.

A bunch of $N_p = 500$ macro-particles is sampled from the 3D parabolic profile $\rho(\boldsymbol{r}) \propto \max(0, 1 - |\boldsymbol{r}|^2/R_b^2)$ ($R_b = 0.55$, matched RMS size $\sigma_0 = R_b/\sqrt{7} = 0.2079$) with an initial centroid offset $(\bar{x}_0, \bar{y}_0) = (0.015, -0.012)$, an 18% quadrupole envelope mismatch ($s_{x,\text{init}} =$

$1.18\sigma_0, s_{y,\text{init}} = 0.856\sigma_0$), and bare betatron tunes $Q_{x0} = 6.23, Q_{y0} = 6.20$, and integrated over $N_{\text{turns}} = 80$ turns ($K_{\text{sub}} = 4$ Leapfrog substeps per turn). At every substep, the physical space-charge disturbance acting on particle p is computed independently via direct $O(N_p^2)$ 3D pairwise Plummer-Coulomb summation in the co-moving frame:

$$\boldsymbol{E}_{\text{true}}(\boldsymbol{r}_p) = \frac{1}{N_p} \sum_{q \neq p}^{N_p} \frac{\boldsymbol{r}_p - \boldsymbol{r}_q}{\left(|\boldsymbol{r}_p - \boldsymbol{r}_q|^2 + h_p^2\right)^{3/2}}, \qquad h_p = 1.6\, R_b N_p^{-1/3}, \tag{27}$$

with effective laboratory-frame perveance $\kappa_{\text{sc}} = 4\pi|\Delta Q_{\text{sc,eff}}|/\langle \partial_x E_x^{\text{true}} \rangle_{\text{core}}$ (incorporating the relativistic $(\beta\gamma^2)^{-1}$ factor via Eq. (1)) calibrated at $t = 0$ to $\Delta Q_{\text{sc,eff}} = \lambda_{\text{sc}} \Delta Q_{\text{sc}} = -0.1817$ ($\lambda_{\text{sc}} = 0.80$).

Crucially, the feedback controller never accesses $\boldsymbol{E}_{\text{true}}(\boldsymbol{r}_p)$ and never performs any particle-level mean subtraction (`np.mean`). Instead, for each actuator family (**B**: 11-term dipole/quadrupole/octupole subset; **C**: full 34-term basis; **D**: full 34-term basis plus turn-by-turn $\Delta s_{200}, \Delta s_{020}$ quadrupole envelope damper), we compare three single-variable configurations:

1. **Laboratory Frame, No ZSF (`B-Lab`, `C-Lab`, `D-Lab`)**: Moments $\boldsymbol{s} = \frac{1}{N_p} \sum_p \boldsymbol{m}(\tilde{\boldsymbol{r}}_p)$ are evaluated in uncentered laboratory coordinates $\boldsymbol{r}_p$ (clipped at $|\tilde{\boldsymbol{r}}_p| \leq 1.15 R_b$), and the aperture-tapered field $\boldsymbol{E}_{\text{model}}(\boldsymbol{r}_p) = -g_{\text{kern}} w_{\text{core}}(r_p) \tilde{\boldsymbol{\Omega}}_p^T \boldsymbol{v}_E^{\text{prior}}$ is applied without ZSF projection, where $w_{\text{core}}(r_p) = \exp[-\frac{1}{2} \max(0, (r_p - R_b)/(0.25 R_b))^2]$.

2. **Centroid Frame, No ZSF (`B-CF`, `C-CF`, `D-CF`)**: Evaluation coordinates are shifted to the instantaneous beam centroid $\boldsymbol{r}_{c,p} = \boldsymbol{r}_p - \bar{\boldsymbol{r}}$ (aligning the bunch with the origin-centered reference prior $\boldsymbol{\mu}$), applying a spatially uniform dipole kick $-g_{\text{kern}} v_{E,l_{x_j}}^{\text{prior}}$ plus aperture-windowed higher-order multipoles ($p + q + l \geq 2$) with `use_zsf=False`.

3. **Centroid Frame + Windowed-Moment ZSF Projection (`B-CF-ZSF`, `C-CF-ZSF`, `D-CF-ZSF`)**: In the exact same centroid frame $\boldsymbol{r}_{c,p}$, Corollary 4 is applied to the unweighted moment vector $\boldsymbol{s}$ (so the polynomial self-force $\|\langle \boldsymbol{E}_{\text{poly}} \rangle\|_\infty$ vanishes to $\sim 10^{-16}$) and to the aperture-windowed 35-moment vector $\boldsymbol{s}_{\text{win}} = \frac{1}{N_p} \sum_{p=1}^{N_p} w_{\text{core}}(r_{c,p}) \boldsymbol{m}(\tilde{\boldsymbol{r}}_{c,p})$ with $s_{\text{win},000} = 1$ for the spatially uniform dipole kicker:

$$E_{x_j,\text{model}}^{\text{ZSF}}(\boldsymbol{r}_{c,p}) = -g_{\text{kern}} \left[v_{E,l_{x_j}}^{\text{ZSF}}(\boldsymbol{s}_{\text{win}}) + w_{\text{core}}(r_{c,p}) \sum_{\substack{l=1 \\ l \neq l_{x_j}}}^{34} v_{E,l}^{\text{prior}} \frac{\partial m_l(\tilde{\boldsymbol{r}}_{c,p})}{\partial x_j} \right]. \tag{28}$$

   By Eq. (23), averaging Eq. (28) over all N_p particles yields $\frac{1}{N_p} \sum_{p=1}^{N_p} \boldsymbol{E}_{\text{model}}^{\text{ZSF}}(\boldsymbol{r}_{c,p}) = \boldsymbol{0}$ to machine precision (1.56–2.42×10^{-16}) *purely from the 35 moments* $\boldsymbol{s}_{\text{win}}$ *and Corollary 4, with zero particle-level mean subtraction.*

Note that while the underlying betatron lattice and pairwise Coulomb kicks are integrated via a symplectic Leapfrog map, the radial aperture window w_{core}, coordinate clipping, and Mode D rate damper make the closed-loop feedback map an aperture-regularized non-Hamiltonian control law rather than a strictly symplectic map. High-action halo fraction is quantified via the normalized transverse betatron action (in dimensionless canonical phase-space units where the linear betatron transfer map is a pure rotation of radius σ_0):

$$\frac{J_\perp}{\sigma_0^2} = \frac{(x - \bar{x})^2 + (p_x - \bar{p}_x)^2 + (y - \bar{y})^2 + (p_y - \bar{p}_y)^2}{2\sigma_0^2} > 3.5, \tag{29}$$

where for $N_p = 500$ macro-particles, a 10% halo fraction corresponds to 50 particles (0.2% resolution per particle). We also track the spurious RMS center-of-mass trajectory distortion

Table 1: De-Confounded Ten-Mode Ablation Study of Idealized Multi-Turn 2D-in-3D Pairwise-Coulomb Leapfrog Tracking Separating Laboratory-Frame vs. Centroid-Frame Alignment and Corollary 4 ZSF Dipolar Projection ($N_{\text{turns}} = 80, N_p = 500, \Delta Q_{\text{sc,eff}} = -0.1817$, Identical Initial Conditions, Seed 42)

Ablation Mode	Frame & ZSF	Max Poly Self-Force	Max Applied Self-Force	Centroid Orbit Err ϵ_{cm}	RMS Emittance Growth	RMS Env. Mismatch	Halo ($J_\perp > 3.5\sigma_0^2$)
Mode A: Open-Loop	Uncompensated	0.00	0.00	0.00	+40.65%	16.60%	29.60%
Mode B-Lab: 11-Term Multipole	Lab, No ZSF	0.4170	0.3931	3.39×10^{-2}	+27.51%	8.04%	16.40%
Mode B-CF: 11-Term Multipole	Centroid, No ZSF	0.3592	0.3549	4.87×10^{-2}	+17.29%	7.06%	16.40%
Mode B-CF-ZSF: 11-Term Multipole	**Centroid + ZSF**	$\mathbf{1.49 \times 10^{-16}}$	$\mathbf{1.56 \times 10^{-16}}$	$\mathbf{4.85 \times 10^{-17}}$	**+17.29%**	**7.06%**	**16.40%**
Mode C-Lab: Full 34-Basis	Lab, No ZSF	0.1833	0.1731	3.66×10^{-2}	+16.44%	3.74%	17.60%
Mode C-CF: Full 34-Basis	Centroid, No ZSF	0.1680	0.1522	6.00×10^{-2}	+13.45%	3.98%	17.00%
Mode C-CF-ZSF: Full 34-Basis	**Centroid + ZSF**	$\mathbf{3.84 \times 10^{-16}}$	$\mathbf{2.13 \times 10^{-16}}$	$\mathbf{4.17 \times 10^{-17}}$	**+13.45%**	**3.98%**	**17.00%**
Mode D-Lab: Full 34 + $\dot{s}_{200}$ Damper	Lab, No ZSF	0.1848	0.1658	3.38×10^{-2}	+7.00%	4.11%	12.20%
Mode D-CF: Full 34 + $\dot{s}_{200}$ Damper	Centroid, No ZSF	0.1506	0.1474	8.12×10^{-2}	+6.40%	4.22%	13.40%
Mode D-CF-ZSF: Full 34 + Damper	**Centroid + ZSF**	$\mathbf{3.98 \times 10^{-16}}$	$\mathbf{2.42 \times 10^{-16}}$	$\mathbf{5.66 \times 10^{-17}}$	+6.40% (6.4× **lower**)	4.22% (3.9× **lower**)	13.40% (2.2× **lower**)

$\epsilon_{\text{cm}} = [\frac{1}{N_{\text{turns}}+1} \sum_{t=0}^{N_{\text{turns}}} |\bar{\boldsymbol{r}}_\perp(t) - \bar{\boldsymbol{r}}_{\perp,\text{open}}(t)|^2]^{1/2}$ relative to the true physical Open-Loop centroid orbit.

Table 1 cleanly resolves how centroid-frame alignment and Corollary 4's ZSF dipolar projection contribute to closed-loop performance:

1. **Physical Effect of Centroid-Frame Alignment (`Lab` → `CF`)**: Because the reference prior $\boldsymbol{\mu}$ is centered at the origin, evaluating the 4th-order polynomial in the moving beam centroid frame $\boldsymbol{r}_{c,p} = \boldsymbol{r}_p - \bar{\boldsymbol{r}}$ eliminates off-axis sextupolar and octupolar feed-down ($(x_c + \bar{x})^3$), reducing internal RMS emittance growth from +27.51% → +17.29% (**Mode B**), +16.44% → +13.45% (**Mode C**), and +7.00% → +6.40% (**Mode D**). (Conversely, applying ZSF in the uncentered laboratory frame removes the accidental origin-restoring dipole damping of the prior while leaving the polynomial off-center, which amplifies off-axis feed-down and confirms why centroid alignment is essential.)

2. **Physical Effect of Corollary 4 ZSF Dipolar Projection (`CF` → `CF-ZSF`)**: Within the centroid frame, Corollary 4 modifies *only* the three spatially uniform dipolar coefficients $(v_{E,x}, v_{E,y}, v_{E,z})$ while leaving all 31 higher-order multipole gradients ($p + q + l \geq 2$) untouched. In a linear betatron lattice, a spatially uniform dipole kick acts purely on the center-of-mass coordinates $(\bar{\boldsymbol{r}}, \bar{\boldsymbol{p}})$ and leaves central internal moments ($\Delta\epsilon/\epsilon_0$, RMS envelope mismatch, and halo fraction) strictly invariant (+17.29%, +13.45%, +6.40%). Simultaneously, ZSF projection eliminates both the polynomial self-force (0.15–0.36 →∼ 10^{-16}) and the applied particle self-force (0.15–0.35 →∼ 10^{-16}), reducing spurious center-of-mass trajectory distortion ϵ_{cm} from 0.049–0.081 to machine zero (4.17–5.66×10^{-17}).

To verify statistical robustness across independent macro-particle realizations, repeating the 80-turn tracking across four random seeds (seed $\in \{42, 1, 7, 99\}$) yields the following per-seed emittance growths, means, and sample standard deviations (ddof = 1, with population standard deviations ddof = 0 in parentheses) across all modes:

- **Mode A (Open-Loop)**: $\{+40.65\%, +53.20\%, +60.85\%, +46.37\%\} \rightarrow +50.27\% \pm 8.72\%$ (pop. SD = 7.55%).
- **Mode B-Lab**: $\{+27.51\%, +19.68\%, +21.32\%, +19.03\%\} \rightarrow +21.88\% \pm 3.87\%$ (pop. SD = 3.35%).
- **Modes B-CF & B-CF-ZSF**: $\{+17.29\%, +16.39\%, +14.86\%, +17.54\%\} \rightarrow +16.52\% \pm 1.21\%$ (pop. SD = 1.05%).
- **Mode C-Lab**: $\{+16.44\%, +17.53\%, +22.03\%, +14.26\%\} \rightarrow +17.56\% \pm 3.27\%$ (pop. SD = 2.83%).
- **Modes C-CF & C-CF-ZSF**: $\{+13.45\%, +22.34\%, +16.46\%, +16.29\%\} \rightarrow +17.13\% \pm 3.73\%$ (pop. SD = 3.23%).

- **Mode D-Lab**: $\{+7.00\%, +8.76\%, +10.86\%, +3.16\%\} \rightarrow +7.44\% \pm 3.26\%$ (pop. SD = 2.83%).
- **Modes D-CF & D-CF-ZSF**: $\{+6.40\%, +10.11\%, +5.05\%, +7.60\%\} \rightarrow +7.29\% \pm 2.15\%$ (pop. SD = 1.86%).

Furthermore, a macro-particle particle-resolution sensitivity check across $N_p \in \{500, 1000, 2000\}$ (Seed 42) suggests that the relative hierarchy of the feedback modes is broadly preserved as collisional discretization noise decreases: Open-Loop **Mode A** grows by (+40.65%, +38.04%, +48.55%), **Mode B-CF-ZSF** yields (+17.29%, +14.74%, +15.29%), **Mode C-CF-ZSF** yields (+13.45%, +13.64%, +15.40%), and **Mode D-CF-ZSF** achieves (+6.40%, +7.05%, +8.99%), corresponding to a 5.4–6.4× reduction relative to Open-Loop across all N_p.

Table 2: Comparison of 3D Space-Charge Compute Kernels (T_{compute}) for High-Intensity Ring Control (Literature Ranges vs. Measured Benchmark Latency)

Solver / Method	**Resolution** ($N_p/G/M$)	**Computational Scaling**	**Kernel Latency** T_{compute} **(CPU)**	**5-Turn Compute Budget** ($\leq 5\,\mu$s)
Direct Pairwise Coulomb (Measured / Lit.)	$N = 10^3$–10^6	$O(N^2)$	1.60 ms–10^2 s	Impossible
3D Grid FFT / Multigrid PIC (Literature†)	$N_p = 10^5, G = 64^3$	$O(N_p + G \log G)$	5.0–50 ms	Unfeasible
Lie-Algebraic Map Truncation (Literature†)	Order 4–6 (10^4 particles)	Symplectic Map Evaluation	2.0–12.0 ms	Unfeasible
Moment Kernel $\boldsymbol{K}, \boldsymbol{K}_E$ **(Measured)**	**Reduced** ($N = 10^3 \rightarrow M = 35$)	**Closed-Form** $O(M_E M)$ (1,190 **MACs**)	**1.6–5.8** $\boldsymbol{\mu}$**s (CPU)**	**Feasible** (T_{compute})
Moment Kernel + ZSF (Measured / FPGA)	**Reduced** ($M = 35, M_E = 34$)	$O(M_E M + M_E)$ (1,247 **MACs**)	**5.4–20** $\boldsymbol{\mu}$**s (Python) / < 1.0** $\boldsymbol{\mu}$**s (FPGA est.)**	**Feasible on C / FPGA**

†Illustrative order-of-magnitude single-step execution times reported across typical CPU hardware [10,11]; Moment Kernel latencies are measured directly in `benchmark_accelerator_space_charge.py`.

—

6 Historical Perspective: Complementary Roles of Stochastic and Moment Control

It is instructive to contrast the Moment Approach with Simon van der Meer's stochastic cooling [8,9]:

1. **Stochastic Cooling**: Samples incoherent Schottky fluctuations of individual particles at high bandwidth to slowly reduce single-particle phase-space entropy over 10^4–10^7 turns in low-bunch-current storage rings.
2. **Moment-Space Multipole Control**: Evaluates coherent macroscopic spatial moments ($M = 35$) in microseconds to compensate collective internal Coulomb defocusing and envelope mismatch oscillations over a few turns (1–100 turns) in space-charge-dominated high-intensity rings.

—

7 Conclusion

We have presented the mathematical formulation and open-source numerical verification of the 1994 Los Alamos Moment Approach for 3D space-charge potential and electric-field reduction in high-intensity hadron rings, augmented by an exact moment-space Zero-Self-Force (ZSF) dipolar offset projection. By regularizing the underdetermined mesh inversion ($\text{rank}(\boldsymbol{\Omega}) = 35$) with a normalized physical prior ($\sum_{i=1}^{N} \mu_i = 1$) and linearizing the weighted least-squares projection around $\boldsymbol{w} = \boldsymbol{\mu}$, the mesh dimension N is algebraically eliminated offline into fixed 35×35 potential and 34×35 electric-field matrices with a compute-kernel latency $T_{\text{compute}} \sim 1.6$–$5.8\,\mu$s on a single CPU core ($\sim 5.4$–$20\,\mu$s with Python-level ZSF indexing, consistently exceeding the $4.73\,\mu$s 5-turn window of the SNS ring in interpreted Python; and $< 1.0\,\mu$s projected on compiled C/FPGA hardware for 1,247 total MACs). These figures represent T_{compute} alone; end-to-end feedback latency $T_{\text{total}} = T_{\text{diag}} + T_{\text{obs}} + T_{\text{compute}} + T_{\text{comm}} + T_{\text{act}}$ was not demonstrated in this work. Combining the prior-weighted operator with the $O(M_E)$ moment-space ZSF dipolar

projection guarantees exact zero net self-force ($\|\langle \boldsymbol{E}^{\mathrm{ZSF}}_{\mathrm{approx}} \rangle_{\boldsymbol{a}}\|_\infty = 4.02 \times 10^{-16}$) on perturbed distributions without accessing the N-point mesh. The physical field approximation accuracy of the 4th-order polynomial basis within the high-density bunch core ($r \leq 0.45$) is $R^2 = 0.8878$ ($\epsilon_E = 33.50\%$ RMS) for the ZSF-projected operator, reflecting the inherent constraint of a fixed-prior perturbative linearization at degree $P \leq 4$; higher-order terms and state-adaptive weights would be required to reduce this residual. Verified in a reproducible de-confounded ten-mode, multi-seed, and multi-N_p 2D-in-3D pairwise-Coulomb Leapfrog tracking study from identical initial conditions, closed-loop moment-operator feedback demonstrates reduced-order T_{compute} feasibility and provides a mathematically transparent foundation for fast envelope mismatch damping and space-charge compensation.

—

Code and Data Availability

The open-source Python reference engine and reproducible verification suite are publicly available on GitHub at:

https://github.com/hyoahn/accelerator-space-charge

The repository includes:

- **Core Accelerator Engine**: `Core/accelerator_space_charge_engine.py` (`AcceleratorSpaceChargeEngine`, implementing the 35×35 potential reduction $\boldsymbol{v} = \boldsymbol{K}\boldsymbol{s} + \boldsymbol{v}_0$, the 34×35 direct 3D electric-field reduction $\boldsymbol{v}_E^{\mathrm{prior}} = \boldsymbol{K}_E\boldsymbol{s} + \boldsymbol{v}_{E,0}$, the moment-space ZSF dipolar offset projection $\boldsymbol{v}_E^{\mathrm{ZSF}}$, three-operator core fidelity analysis, and the de-confounded 10-mode non-circular pairwise-Coulomb Leapfrog tracking ablation and multi-seed robustness suite).
- **Verification & Benchmark Suite**: `Test/benchmark_accelerator_space_charge.py` (standalone-import compatible, reproducing all numerical residuals, three-operator R^2 checks, ZSF dipolar projection checks, latency benchmarks, and 10-mode de-confounded ablation metrics reported in Section 5, Fig. 1, Table 1, and Table 2).

—

Historical Note and Acknowledgments

The mathematical formulations of the normalized-prior charge inversion ('report1.pdf'), the polynomial potential reduction ('poly-phi.pdf'), and the direct 3D vector electric-field reduction ('poly-e.pdf'), including the algebraic elimination of the spatial grid dimension N (the fall-out of N) and the weighted moment and force identities, were originally derived by the author in 1994 as unpublished research notes at Los Alamos National Laboratory. The moment-space Zero-Self-Force (ZSF) dipolar offset projection (Corollary 4) and the de-confounded multi-turn pairwise-Coulomb Leapfrog ablation benchmarks were developed for the present 2026 study. The author acknowledges early collaborative discussions with Walter P. Lysenko in 1993–1994 at Los Alamos National Laboratory and the early conceptual ideas of the late Paul J. Channell regarding underdetermined moment models for the BEDLAM simulation code.

—

References

[1] M. Reiser, *Theory and Design of Charged Particle Beams*, John Wiley & Sons, 2nd ed., 2008.

[2] R. Baartman, "Betatron resonances with space charge," *AIP Conference Proceedings*, vol. 448, no. 1, pp. 56–72, 1998.

[3] A. Abada, et al. (FCC Collaboration), "FCC-hh: The Hadron Collider: Future Circular Collider Conceptual Design Report Volume 3," *Eur. Phys. J. Spec. Top.*, vol. 228, no. 4, pp. 755–1107, 2019.

[4] S. Henderson, et al., "The Spallation Neutron Source accelerator system design," *Nucl. Instrum. Methods Phys. Res. A*, vol. 763, pp. 610–673, 2014.

[5] E. Benedetto, et al., "Space charge effects and mitigation in the CERN PS Booster in view of the LIU upgrade," in *Proc. IPAC'15*, Richmond, VA, USA, pp. 3899–3902, 2015.

[6] P. Spiller and G. Franchetti, "The FAIR accelerator project at GSI," *Nucl. Instrum. Methods Phys. Res. A*, vol. 561, no. 2, pp. 305–312, 2006.

[7] A. V. Fedotov, I. Hofmann, R. L. Gluckstern, and H. Okamoto, "Parametric collective resonances and space-charge limit in high-intensity rings," *Phys. Rev. ST Accel. Beams*, vol. 6, p. 094201, 2003.

[8] S. van der Meer, "Stochastic damping of betatron oscillations in the ISR," CERN Report CERN/ISR-PO/72-31, 1972.

[9] S. van der Meer, "Stochastic cooling and the accumulation of antiprotons (Nobel Lecture)," *Rev. Mod. Phys.*, vol. 57, no. 3, pp. 689–697, 1985.

[10] C. K. Birdsall and A. B. Langdon, *Plasma Physics via Computer Simulation*, CRC Press, 2018.

[11] J. Qiang, R. D. Ryne, S. Habib, and V. Decyk, "An object-oriented parallel Particle-in-Cell code for beam dynamics simulation in linear accelerators," *J. Comput. Phys.*, vol. 163, no. 2, pp. 434–451, 2000.

[12] P. J. Channell, "The moment approach to space charge in particle beams," Los Alamos National Laboratory Report, 1993.

[13] H. E. Ahn and W. P. Lysenko, "Preliminary Results of Space-Charge Model Study," Los Alamos National Laboratory Internal Memorandum (Unpublished), April 4, 1994.