Closed-Form Moment Reduction of 3D Space-Charge Fields for Fast Feedback in High-Intensity Hadron Rings

Hyo Eun Ahn*
Independent Researcher

November 2, 2026

Abstract

In high-intensity hadron accumulator rings, rapid-cycling synchrotrons, and low-to-medium energy injector boosters—such as the Los Alamos Proton Storage Ring (PSR), the Spallation Neutron Source (SNS) accumulator ring, the CERN Proton Synchrotron Booster (PSB), and FAIR SIS18—intense self-induced Coulomb repulsion inside dense particle bunches ($\gamma \sim 1.15$–3.1) drives large incoherent space-charge tune shifts ($\Delta Q_{\mathrm{sc}} \sim -0.1$ to -0.4), resulting in envelope mismatch resonances, transverse emittance growth, and halo generation. While active damping of dipolar centroid oscillations has been standard practice since Simon van der Meer's stochastic cooling, real-time active compensation of higher-order space-charge fields remains computationally challenging due to the millisecond-to-second latency of direct $O(N^2)$ Coulomb summation and grid-based $O(N_p + G \log G)$ Particle-in-Cell (PIC) solvers relative to the sub-microsecond to microsecond revolution period ($T_{\mathrm{rev}} \approx 0.36$–$1.0\,\mu$s, or $\leq 5\,\mu$s for a 5-turn envelope control window). In this paper, building on the author's 1994 Los Alamos space-charge moment formulations, we establish a closed-form reduced-order moment framework that evaluates 3D interior space-charge potential and electric-field polynomial coefficients with a compute-kernel latency of ~ 1.6–$4.6\,\mu$s on a single CPU core. By expanding the bunch distribution into $M = 35$ spatial moments (all 3D monomials up to 4th order) and regularizing the underdetermined N-point mesh inversion ($\mathrm{rank}(\boldsymbol{\Omega}) = 35$) with a normalized physical reference prior ($\sum_{i=1}^{N} \mu_i = 1, \mu_i > 0$), strict pointwise positivity ($\boldsymbol{a} \equiv \boldsymbol{\mu} > \mathbf{0}$) is guaranteed at the reference equilibrium and fringe negative-charge overshoot is suppressed by $3.04\times$ ($0.2669 \to 0.0876$) under finite moment perturbations compared to Channell's uniform prior. Using reference-prior linearization ($\boldsymbol{w} = \boldsymbol{\mu}$), we demonstrate the exact algebraic "fall-out" of the mesh dimension N, collapsing the prior-weighted polynomial projection of the Plummer-softened Coulomb kernel into precomputed 35×35 potential ($\boldsymbol{v} = \boldsymbol{K}\boldsymbol{s} + \boldsymbol{v}_0$) and 34×35 electric-field ($\boldsymbol{v}_E^{\mathrm{prior}} = \boldsymbol{K}_E\boldsymbol{s} + \boldsymbol{v}_{E,0}$) matrices. Furthermore, because the uncorrected prior-weighted operator guarantees $\langle \boldsymbol{E}_{\mathrm{approx}} \rangle_{\boldsymbol{\mu}} = \langle \boldsymbol{E}_{\mathrm{exact}} \rangle_{\boldsymbol{\mu}}$ rather than vanishing net force on a perturbed state $\boldsymbol{a} = \boldsymbol{\mu} + \delta\boldsymbol{a}$ (where $\boldsymbol{\mu}^T \boldsymbol{G} \delta\boldsymbol{a} \neq \mathbf{0}$), we introduce an exact $O(M_E)$ moment-space Zero-Self-Force (ZSF) dipolar gauge projection that restores $\langle \boldsymbol{E}_{\mathrm{approx}}^{\mathrm{ZSF}} \rangle_{\boldsymbol{a}} = \mathbf{0}$ to machine precision (4.02×10^{-16}) without accessing the N-point mesh. In an idealized zero-latency multi-turn 2D transverse Leapfrog tracking and four-mode ablation study driven by self-consistent $O(N_p^2)$ 3D pairwise Plummer-Coulomb forces from identical initial conditions, moment-operator feedback reduces RMS envelope mismatch excursion by $4.4\times$ ($16.60\% \to 3.74\%$) and RMS emittance growth from $+40.65\%$ to $+16.44\%$ ($+7.00\%$ when combined with a 2nd-moment rate damper).

Keywords: Space-Charge Field Reduction, High-Intensity Hadron Rings, Injector Boosters, Moment Approach, Reduced-Order Modeling, Fast Multipole Feedback.

—

*Email: ahn@aya.yale.edu

1 Introduction

The ultimate beam brightness (N_b/ϵ_n) delivered by high-power proton drivers and hadron collider injector chains is fundamentally limited by collective space-charge forces at low and intermediate energies [1, 2]. In the laboratory frame, the self-generated magnetic force of a relativistic bunch partially cancels the electrostatic repulsion by a factor of $1-\beta^2=\gamma^{-2}$, so that the incoherent Laslett space-charge betatron tune shift scales inversely with $\beta\gamma^2$:

$$\Delta Q_{\mathrm{sc}} = -\frac{r_p N_b F_g}{2\pi\epsilon_n \beta\gamma^2 B_f}, \tag{1}$$

where $r_p = q^2/(4\pi\epsilon_0 m_p c^2)$ is the classical proton radius, N_b is the bunch intensity, ϵ_n is the normalized transverse RMS emittance, β and γ are the relativistic Lorentz factors, B_f is the bunching factor, and F_g is the transverse form factor. Consequently, while space-charge tune depression is negligible ($\Delta Q_{\mathrm{sc}} \sim 10^{-8}$–$10^{-9}$) at multi-TeV collider top energies [3], it constitutes the primary intensity bottleneck in low-to-medium energy accumulator rings and injector boosters ($\gamma \sim 1.15$–3.1). Specifically, in machines such as the Los Alamos Proton Storage Ring (PSR, $E_k = 800\,\mathrm{MeV}$, $\gamma = 1.853$, circumference $C = 90.2\,\mathrm{m}$, revolution period $T_{\mathrm{rev}} = 0.358\,\mu\mathrm{s}$), the Spallation Neutron Source (SNS) accumulator ring ($E_k = 1.0\,\mathrm{GeV}$, $\gamma = 2.066$, $C = 248.0\,\mathrm{m}$, $T_{\mathrm{rev}} = 0.945\,\mu\mathrm{s}$) [4], the CERN Proton Synchrotron Booster (PSB, $E_k = 160\,\mathrm{MeV}$–$2.0\,\mathrm{GeV}$, $T_{\mathrm{rev}} = 0.57$–$1.01\,\mu\mathrm{s}$) [5], and the FAIR SIS18 synchrotron [6], the incoherent tune depression reaches $|\Delta Q_{\mathrm{sc}}| \approx 0.1$–$0.4$.

When this large tune shift couples with injection envelope mismatch or pushes particles across low-order structure resonances ($m_x Q_x + m_y Q_y = k$), parametric collective resonances drive particles out of the high-density bunch core into an extended high-action halo [7], causing beam loss and radioactivation of the ring aperture.

Historically, active feedback in circular accelerators has been restricted to dipolar (center-of-mass) centroid damping. Simon van der Meer's stochastic cooling [8, 9], awarded the 1984 Nobel Prize in Physics, damps incoherent phase-space fluctuations over 10^4–10^7 turns. By contrast, coherent envelope oscillations and space-charge-induced mismatch in compact high-intensity rings evolve over a few revolution turns ($T_{\mathrm{rev}} \approx 0.36$–$1.0\,\mu\mathrm{s}$, corresponding to a 5-turn control window of $\leq 5\,\mu\mathrm{s}$). Conventional self-consistent space-charge solvers cannot execute within this microsecond compute budget: direct pairwise Coulomb summation over N mesh or macro-particle nodes scales as $O(N^2)$ ($\sim 1.6\,\mathrm{ms}$ even for $N = 10^3$), while grid-based Fast Fourier Transform (FFT) or multigrid Particle-in-Cell (PIC) codes [10, 11] scale as $O(N_p + G\log G)$ for N_p macro-particles and G grid cells, requiring milliseconds to seconds per step.

In this work, we present the mathematical formalization and open-source numerical verification of a closed-form reduced-order Moment Approach to 3D space-charge field estimation and fast feedback control. It is important to distinguish the two components of this contribution:

1. **1994 LANL Reduced-Order Moment Formulation (Sections 2–4)**: Originally motivated by Channell's moment dynamics [12] and initiated during 1994 Los Alamos National Laboratory (LANL) studies for the BEDLAM code [13], the author derived the normalized-prior affine charge inversion ('report1.pdf'), the weighted polynomial potential reduction ('poly-phi.pdf'), and the direct 3D vector electric-field reduction ('poly-e.pdf'), showing that the mesh dimension N is algebraically eliminated offline into fixed 35×35 and 34×35 matrices.

2. **2026 Moment-Space Zero-Self-Force (ZSF) Gauge Projection and Closed-Loop Feasibility Study (Sections 4.2–5)**: Extending the 1994 operator to modern real-time feedback, we resolve the non-zero self-force residual that arises when applying a fixed prior-weighted operator ($\boldsymbol{w} = \boldsymbol{\mu}$) to a perturbed beam distribution ($\boldsymbol{a} \neq \boldsymbol{\mu}$) via an exact $O(M_E)$ moment-space dipolar gauge projection, and evaluate the resulting operator in

a reproducible multi-turn 2D-in-3D pairwise-Coulomb Leapfrog tracking and ablation benchmark.

—

2 Underdetermined Charge Inversion and Prior Regularization

2.1 Moment Representation in 3D Real Space

In the co-moving rest frame of the bunch, let the spatial domain $\mathcal{V} \subset \mathbb{R}^3$ be discretized into N mesh points $\boldsymbol{r}_i = (x_i, y_i, z_i)$ $(i = 1, \ldots, N)$. Let a_i denote the normalized point-charge weight at mesh node $\boldsymbol{r}_i$, with $\boldsymbol{a} = [a_1, \ldots, a_N]^T \in \mathbb{R}^N$.

The macroscopic geometry of the bunch is characterized by M spatial moments up to polynomial order $P = 4$:

$$s_k = \langle m_k(\boldsymbol{r}) \rangle_{\boldsymbol{a}} = \sum_{i=1}^{N} m_k(\boldsymbol{r}_i) a_i, \quad k = 1, \ldots, M, \tag{2}$$

where $m_k(\boldsymbol{r}) = x^{p_k} y^{q_k} z^{l_k}$ $(p_k + q_k + l_k \leq 4)$ are the $M = \binom{4+3}{3} = 35$ distinct 3D monomial basis functions:

$$\boldsymbol{m}(\boldsymbol{r}) = \left[1, x, y, z, x^2, xy, xz, y^2, yz, z^2, x^3, \ldots, z^4\right]^T. \tag{3}$$

Defining the moment evaluation matrix $\boldsymbol{\Omega} \in \mathbb{R}^{M \times N}$ with entries $\Omega_{ki} = m_k(\boldsymbol{r}_i)$, Eq. (2) takes the linear form:

$$\boldsymbol{s} = \boldsymbol{\Omega} \boldsymbol{a}. \tag{4}$$

When $N \gg M$ (typically $N \sim 10^3$–10^5 while $M = 35$) and the mesh points are non-coplanar on algebraic surfaces of degree ≤ 4 such that $\operatorname{rank}(\boldsymbol{\Omega}) = M = 35$, Eq. (4) is an underdetermined linear inverse problem with an $(N - M)$-dimensional nullspace.

2.2 Constrained Variational Inversion and Fringe Negativity Suppression

To select a unique, smooth mesh charge assignment $\boldsymbol{a}$ without iterative non-negative quadratic programming (which would violate microsecond real-time constraints), we minimize the constrained quadratic functional [13]:

$$\mathcal{J}(\boldsymbol{a}, \boldsymbol{\lambda}) = \frac{1}{2} \sum_{i=1}^{N} (a_i - \mu_i)^2 + \boldsymbol{\lambda}^T (\boldsymbol{s} - \boldsymbol{\Omega} \boldsymbol{a}), \tag{5}$$

where $\boldsymbol{\mu} = [\mu_1, \ldots, \mu_N]^T \in \mathbb{R}^N$ is a fixed reference prior distribution and $\boldsymbol{\lambda} \in \mathbb{R}^M$ is the Lagrange multiplier vector.

In early tests for the BEDLAM code at LANL, setting $\boldsymbol{\mu} = \mathbf{0}$ yielded the minimum-L_2-norm pseudoinverse solution $\boldsymbol{a}_{\min} = \boldsymbol{\Omega}^T (\boldsymbol{\Omega} \boldsymbol{\Omega}^T)^{-1} \boldsymbol{s}$, which exhibits large unphysical negative charge oscillations ($a_i < 0$) in the low-density bunch fringe whenever the bunch radius is smaller than the mesh radius ($R_{\text{beam}} < R_{\text{mesh}}$). Paul Channell suggested a uniform reference prior $\mu_i = \text{const}$ [12]. However, because the zeroth-order constant monomial $m_1(\boldsymbol{r}) = 1$ lies in the row space of $\boldsymbol{\Omega}$, the orthogonal projection operator $[\boldsymbol{I}_N - \boldsymbol{\Omega}^T (\boldsymbol{\Omega} \boldsymbol{\Omega}^T)^{-1} \boldsymbol{\Omega}]$ annihilates any constant vector identically ($\|\boldsymbol{a}(\mu = \text{const}) - \boldsymbol{a}(\mu = 0)\|_\infty = 0$), leaving the fringe negativity unchanged.

As shown by the author in the April 1994 LANL study [13], effective non-iterative regularization requires choosing a spatially localized physical reference profile $\boldsymbol{\mu}$ whose second moments match the nominal beam envelope and whose total weight is normalized to unity:

$$\sum_{i=1}^{N} \mu_i = 1, \quad \mu_i > 0 \quad \forall i. \tag{6}$$

Theorem 1 (Closed-Form Affine Charge Reconstruction)**.** *Assume the N-point 3D mesh satisfies the full-row-rank condition* $\operatorname{rank}(\boldsymbol{\Omega}) = M = 35$*, so that $\boldsymbol{\Omega}\boldsymbol{\Omega}^T \in \mathbb{R}^{35\times 35}$ is strictly symmetric positive-definite. Under the normalization constraint (6), the unique global minimizer of $\mathcal{J}(\boldsymbol{a}, \boldsymbol{\lambda})$ subject to $\boldsymbol{\Omega a} = \boldsymbol{s}$ is the affine transformation:*

$$\boldsymbol{a} = \boldsymbol{\mu} + \boldsymbol{\Lambda}(\boldsymbol{s} - \boldsymbol{\Omega\mu}) = \boldsymbol{\Lambda s} + \boldsymbol{c}, \tag{7}$$

where $\boldsymbol{\Lambda} \in \mathbb{R}^{N\times M}$ is the Moore-Penrose right pseudoinverse and $\boldsymbol{c} \in \mathbb{R}^N$ is the prior residual vector:

$$\boldsymbol{\Lambda} = \boldsymbol{\Omega}^T(\boldsymbol{\Omega}\boldsymbol{\Omega}^T)^{-1}, \qquad \boldsymbol{c} = (\boldsymbol{I}_N - \boldsymbol{\Lambda\Omega})\boldsymbol{\mu}. \tag{8}$$

At the reference equilibrium ($\boldsymbol{s} = \boldsymbol{\Omega\mu}$), Eq. (7) yields $\boldsymbol{a} \equiv \boldsymbol{\mu} > \boldsymbol{0}$, guaranteeing strict pointwise positivity across all N mesh points. For perturbed beams ($\boldsymbol{s} \neq \boldsymbol{\Omega\mu}$), Eq. (7) preserves all $M = 35$ moments ($\boldsymbol{\Omega a} = \boldsymbol{s}$) identically, and maintains pointwise positivity ($a_i > 0$) whenever the moment perturbation satisfies the componentwise bound $[\boldsymbol{\Lambda}(\boldsymbol{s} - \boldsymbol{\Omega\mu})]_i > -\mu_i$. Outside this perturbative neighborhood, unconstrained affine inversion permits localized negative fringe weights ($a_i < 0$), but reduces integrated fringe negative-charge overshoot by more than $3\times$ compared to the uniform/zero prior.

Proof. Setting $\nabla_{\boldsymbol{a}}\mathcal{J} = \boldsymbol{0}$ gives $\boldsymbol{a} - \boldsymbol{\mu} - \boldsymbol{\Omega}^T\boldsymbol{\lambda} = \boldsymbol{0}$, or $\boldsymbol{a} = \boldsymbol{\mu} + \boldsymbol{\Omega}^T\boldsymbol{\lambda}$. Multiplying from the left by $\boldsymbol{\Omega}$ and enforcing $\boldsymbol{\Omega a} = \boldsymbol{s}$ yields $\boldsymbol{\Omega\mu} + (\boldsymbol{\Omega}\boldsymbol{\Omega}^T)\boldsymbol{\lambda} = \boldsymbol{s}$. Because $\operatorname{rank}(\boldsymbol{\Omega}) = M$, $\boldsymbol{\Omega}\boldsymbol{\Omega}^T$ is invertible, giving $\boldsymbol{\lambda} = (\boldsymbol{\Omega}\boldsymbol{\Omega}^T)^{-1}(\boldsymbol{s} - \boldsymbol{\Omega\mu})$. Substituting $\boldsymbol{\lambda}$ back into $\boldsymbol{a}$ establishes Eq. (7), and the positivity condition $a_i > 0 \iff [\boldsymbol{\Lambda}(\boldsymbol{s} - \boldsymbol{\Omega\mu})]_i > -\mu_i$ follows immediately. □

—

3 Space-Charge Potential Reduction and the Fall-Out of N

3.1 Interior Polynomial Potential Fitting via Weighted Least Squares

Given a mesh charge distribution $\boldsymbol{a}$ in the bunch rest frame, the regularized discrete electrostatic potential on the mesh is given by the Plummer-softened Coulomb sum:

$$\phi_{\text{exact}}(\boldsymbol{r}_i) = \sum_{j=1}^{N} D_{ij}a_j, \qquad D_{ij} = \begin{cases} \dfrac{1}{\sqrt{|\boldsymbol{r}_i - \boldsymbol{r}_j|^2 + h^2}}, & i \neq j, \\ 0, & i = j, \end{cases} \tag{9}$$

where $h = \alpha_h N^{-1/3}$ is the inter-particle mesh softening scale that prevents singular two-body collisions on the discrete grid and $D_{ii} = 0$ excludes self-interaction.

To represent the interior space-charge potential analytically within the beam core, we expand ϕ in the 4th-order monomial basis:

$$\phi_{\text{approx}}(\boldsymbol{r}) = \sum_{l=1}^{M} v_l m_l(\boldsymbol{r}) = \boldsymbol{m}^T(\boldsymbol{r})\boldsymbol{v}, \tag{10}$$

where $\boldsymbol{v} = [v_1, \ldots, v_M]^T \in \mathbb{R}^M$. Because a 4th-order polynomial grows as r^4 at large radii whereas the exterior Laplace potential decays as r^{-1}, an unweighted fit over the bounding box would be distorted by empty vacuum nodes outside the bunch. We therefore weight the least-squares functional by a positive density measure w_i concentrated in the beam core:

$$\chi^2_{\boldsymbol{w}}(\boldsymbol{v}) = \sum_{i=1}^{N} w_i \left[\phi_{\text{exact}}(\boldsymbol{r}_i) - \phi_{\text{approx}}(\boldsymbol{r}_i)\right]^2. \tag{11}$$

Theorem 2 (Potential Dimension Reduction and Weighted Moment Conservation)**.** *For any weight vector $\boldsymbol{w} \in \mathbb{R}^N$ such that the weighted Gram matrix $\boldsymbol{\Omega}\{\boldsymbol{\Omega w}\}^T \in \mathbb{R}^{M\times M}$ is non-singular (which holds whenever $w_i > 0$ on a rank-M mesh), the normal equations minimizing $\chi^2_{\boldsymbol{w}}(\boldsymbol{v})$ yield the unique coefficient vector:*

$$\boldsymbol{v}(\boldsymbol{w}) = \boldsymbol{\Lambda}'^T(\boldsymbol{w})\boldsymbol{\phi}_{\text{exact}}, \qquad \boldsymbol{\Lambda}'(\boldsymbol{w}) = \{\boldsymbol{\Omega w}\}^T\Big[\boldsymbol{\Omega}\{\boldsymbol{\Omega w}\}^T\Big]^{-1}, \tag{12}$$

where $\{\boldsymbol{\Omega w}\}_{ki} = \Omega_{ki}w_i$. By virtue of the normal equations, the polynomial potential $\boldsymbol{\phi}_{\text{approx}} = \boldsymbol{\Omega}^T\boldsymbol{v}(\boldsymbol{w})$ preserves all $M = 35$ $\boldsymbol{w}$-weighted potential moments identically:

$$\langle m_k\phi_{\text{approx}}\rangle_{\boldsymbol{w}} = \sum_{i=1}^{N} w_i m_k(\boldsymbol{r}_i)\phi_{\text{approx}}(\boldsymbol{r}_i) = \langle m_k\phi_{\text{exact}}\rangle_{\boldsymbol{w}}, \qquad k = 1,\dots,M. \tag{13}$$

Furthermore, adopting the fixed reference-prior weighting $\boldsymbol{w} = \boldsymbol{\mu}$ (first-order perturbative linearization around $\boldsymbol{a} = \boldsymbol{\mu}$) eliminates the mesh dimension N entirely offline, yielding an algebraically exact reduced-order representation of the $\boldsymbol{\mu}$-weighted polynomial projection on the reconstructed charge $\boldsymbol{a} = \boldsymbol{\Lambda s} + \boldsymbol{c}$:

$$\boldsymbol{v} = \boldsymbol{K s} + \boldsymbol{v}_0, \qquad \boldsymbol{K} = \boldsymbol{\Lambda}'^T(\boldsymbol{\mu})\boldsymbol{D}\boldsymbol{\Lambda} \in \mathbb{R}^{M\times M}, \qquad \boldsymbol{v}_0 = \boldsymbol{\Lambda}'^T(\boldsymbol{\mu})\boldsymbol{D}\boldsymbol{c} \in \mathbb{R}^M. \tag{14}$$

Proof. Setting $\partial\chi^2_{\boldsymbol{w}}/\partial v_k = 0$ for $k = 1,\dots,M$ yields the weighted normal equations:

$$\sum_{i=1}^{N} w_i\Omega_{ki}\left(\boldsymbol{\Omega}^T\boldsymbol{v}\right)_i = \sum_{i=1}^{N} w_i\Omega_{ki}\phi_{\text{exact}}(\boldsymbol{r}_i) \iff \{\boldsymbol{\Omega w}\}\boldsymbol{\Omega}^T\boldsymbol{v} = \{\boldsymbol{\Omega w}\}\boldsymbol{\phi}_{\text{exact}}, \tag{15}$$

which establishes the moment conservation identity (13) and, upon left-multiplying by $[\{\boldsymbol{\Omega w}\}\boldsymbol{\Omega}^T]^{-1}$, gives $\boldsymbol{v}(\boldsymbol{w}) = \boldsymbol{\Lambda}'^T(\boldsymbol{w})\boldsymbol{\phi}_{\text{exact}}$. Substituting $\boldsymbol{\phi}_{\text{exact}} = \boldsymbol{D a} = \boldsymbol{D}(\boldsymbol{\Lambda s} + \boldsymbol{c})$ with fixed weight $\boldsymbol{w} = \boldsymbol{\mu}$ factors all N-dimensional matrices into the offline constants $\boldsymbol{K}$ and $\boldsymbol{v}_0$ of Eq. (14). □

Remark on Algebraic Reduction Exactness vs. Physical Field Approximation: It is essential to distinguish two levels of accuracy in Theorem 2: (i) *algebraic exactness of the reduced operator*, meaning that $\boldsymbol{v} = \boldsymbol{K s} + \boldsymbol{v}_0 \in \mathbb{R}^{35}$ reproduces the full N-mesh prior-weighted projection $\boldsymbol{\Lambda}'^T(\boldsymbol{\mu})\boldsymbol{D}(\boldsymbol{\Lambda s} + \boldsymbol{c})$ with zero truncation error (1.14×10^{-13} numerically), and (ii) *physical field approximation*, which involves replacing the true continuum density by the affine reconstruction $\boldsymbol{a}$, truncating the potential basis at degree $P = 4$, and replacing the state-dependent weight $\boldsymbol{w} = \boldsymbol{a}$ by the fixed prior weight $\boldsymbol{w} = \boldsymbol{\mu}$ (introducing an $O(\|\boldsymbol{a} - \boldsymbol{\mu}\|)$ difference relative to the state-weighted fit).

—

4 Direct 3D Vector Electric Field Reduction and Moment-Space Zero-Self-Force Gauge Projection

4.1 Direct Vector Field Least-Squares Formulation

In beam dynamics, betatron tune shifts and multipole kicks depend directly on the electric field $\boldsymbol{E} = -\nabla\phi$. Rather than differentiating a fitted scalar potential (which amplifies high-order boundary oscillations), the author derived a direct 3D vector electric-field least-squares fit in 1994 (‘poly-e.pdf‘).

Using the Plummer-softened field kernels $G_{x_j,ik} = (x_{j,i} - x_{j,k})(|\boldsymbol{r}_i - \boldsymbol{r}_k|^2 + h^2)^{-3/2}$ ($i \neq k$, $G_{x_j,ii} = 0$), the exact discrete electric field components along x_1, x_2, x_3 (x, y, z) are:

$$\boldsymbol{E}^{\text{exact}}_{x_j} = \boldsymbol{G}_{x_j}\boldsymbol{a}, \qquad j = 1,2,3. \tag{16}$$

Because the constant monomial $m_1(\boldsymbol{r}) = 1$ has zero gradient ($\nabla m_1 \equiv \mathbf{0}$), we retain only the $M_E = M - 1 = 34$ non-constant monomials ($1 \le p_l + q_l + l_l \le 4$), denoted by the reduced derivative matrices $\tilde{\boldsymbol{\Omega}}^{x_j} \in \mathbb{R}^{34\times N}$ with entries $\tilde{\Omega}^{x_j}_{li} = \partial m_l(\boldsymbol{r}_i)/\partial x_j$. The polynomial electric field approximation is:

$$\boldsymbol{E}^{\text{approx}}_{x_j} = -\tilde{\boldsymbol{\Omega}}^{x_j T}\boldsymbol{v}_E, \qquad \boldsymbol{v}_E \in \mathbb{R}^{34}. \tag{17}$$

Minimizing the density-weighted 3D vector residual functional:

$$\mathcal{S}_{\boldsymbol{w}}(\boldsymbol{v}_E) = \sum_{i=1}^{N} w_i \sum_{j=1}^{3} \left[E^{\text{exact}}_{x_j}(\boldsymbol{r}_i) - E^{\text{approx}}_{x_j}(\boldsymbol{r}_i)\right]^2, \tag{18}$$

with respect to $\boldsymbol{v}_E$ yields the following theorem.

Theorem 3 (Direct 3D Electric Field Reduction and Weighted Force Identity)**.** *For any weight vector $\boldsymbol{w} \in \mathbb{R}^N$ such that $\sum_{j=1}^{3}\{\tilde{\boldsymbol{\Omega}}^{x_j}\boldsymbol{w}\}\tilde{\boldsymbol{\Omega}}^{x_j T} \in \mathbb{R}^{34\times 34}$ is invertible, the optimal 34-component electric-field coefficient vector minimizing $\mathcal{S}_{\boldsymbol{w}}(\boldsymbol{v}_E)$ is:*

$$\boldsymbol{v}_E(\boldsymbol{w}) = -\left[\sum_{j=1}^{3}\{\tilde{\boldsymbol{\Omega}}^{x_j}\boldsymbol{w}\}\tilde{\boldsymbol{\Omega}}^{x_j T}\right]^{-1}\left[\sum_{j=1}^{3}\{\tilde{\boldsymbol{\Omega}}^{x_j}\boldsymbol{w}\}\boldsymbol{E}^{\text{exact}}_{x_j}\right]. \tag{19}$$

This solution satisfies all $M_E = 34$ $\boldsymbol{w}$-weighted generalized field-moment normal equations identically:

$$\sum_{j=1}^{3}\left\langle \frac{\partial m_k}{\partial x_j} E^{\text{approx}}_{x_j}(\boldsymbol{w})\right\rangle_{\boldsymbol{w}} = \sum_{j=1}^{3}\left\langle \frac{\partial m_k}{\partial x_j} E^{\text{exact}}_{x_j}\right\rangle_{\boldsymbol{w}}, \qquad k = 2,\dots,35. \tag{20}$$

*In particular, evaluating Eq. (20) for the three linear dipolar monomials $m_k \in \{x, y, z\}$ (where $\partial x_i/\partial x_j = \delta_{ij}$) guarantees the $\boldsymbol{w}$-weighted **Average Force Identity**:*

$$\langle \boldsymbol{E}_{\text{approx}}(\boldsymbol{w})\rangle_{\boldsymbol{w}} = \langle \boldsymbol{E}_{\text{exact}}\rangle_{\boldsymbol{w}}. \tag{21}$$

Consequently:

1. ***State-Weighted Fit ($\boldsymbol{w} = \boldsymbol{a}$)**: By the antisymmetry $G_{x_j,ik} = -G_{x_j,ki}$ (Newton's third law), $\langle E^{\text{exact}}_{x_j}\rangle_{\boldsymbol{a}} = \boldsymbol{a}^T\boldsymbol{G}_{x_j}\boldsymbol{a} = 0$, so the state-weighted fit satisfies $\langle \boldsymbol{E}_{\text{approx}}(\boldsymbol{a})\rangle_{\boldsymbol{a}} = \mathbf{0}$ identically.*
2. ***Prior-Weighted Offline Reduction ($\boldsymbol{w} = \boldsymbol{\mu}$)**: Fixing $\boldsymbol{w} = \boldsymbol{\mu}$ collapses Eq. (19) offline into the runtime 34×35 affine operator:*

$$\boldsymbol{v}^{\text{prior}}_E = \boldsymbol{K}_E\boldsymbol{s} + \boldsymbol{v}_{E,0}, \qquad \boldsymbol{K}_E = \boldsymbol{\Lambda}^T_E(\boldsymbol{\mu})\boldsymbol{\Lambda} \in \mathbb{R}^{34\times 35}, \qquad \boldsymbol{v}_{E,0} = \boldsymbol{\Lambda}^T_E(\boldsymbol{\mu})\boldsymbol{c} \in \mathbb{R}^{34}, \tag{22}$$

where $\boldsymbol{\Lambda}^T_E(\boldsymbol{\mu}) = -\left[\sum_{j=1}^{3}\{\tilde{\boldsymbol{\Omega}}^{x_j}\boldsymbol{\mu}\}\tilde{\boldsymbol{\Omega}}^{x_j T}\right]^{-1}\left(\sum_{j=1}^{3}\{\tilde{\boldsymbol{\Omega}}^{x_j}\boldsymbol{\mu}\}\boldsymbol{G}_{x_j}\right) \in \mathbb{R}^{34\times N}$. For a perturbed state $\boldsymbol{a} = \boldsymbol{\mu} + \delta\boldsymbol{a}$, Eq. (22) preserves the $\boldsymbol{\mu}$-weighted field averages $\langle E^{\text{approx}}_{x_j}(\boldsymbol{\mu})\rangle_{\boldsymbol{\mu}} = \langle E^{\text{exact}}_{x_j}\rangle_{\boldsymbol{\mu}} = \boldsymbol{\mu}^T\boldsymbol{G}_{x_j}\delta\boldsymbol{a} \neq 0$, which generally leaves a non-zero residual net force $\langle \boldsymbol{E}_{\text{approx}}(\boldsymbol{\mu})\rangle_{\boldsymbol{a}} \neq \mathbf{0}$ on the actual distribution $\boldsymbol{a}$.

4.2 Exact Moment-Space Zero-Self-Force (ZSF) Dipolar Gauge Projection

As highlighted in Theorem 3, because $\boldsymbol{w} = \boldsymbol{\mu}$ is fixed offline, the uncorrected prior-weighted operator $\boldsymbol{v}^{\text{prior}}_E$ does not automatically vanish when averaged over a perturbed charge state $\boldsymbol{a} \neq \boldsymbol{\mu}$: both $\boldsymbol{\mu}^T\boldsymbol{G}_{x_j}\delta\boldsymbol{a} \neq 0$ and the weighting mismatch $(\boldsymbol{a} - \boldsymbol{\mu})^T\boldsymbol{E}^{\text{approx}}_{x_j}(\boldsymbol{\mu}) \neq 0$ contribute a spurious dipolar self-force $\langle \boldsymbol{E}_{\text{approx}}(\boldsymbol{\mu})\rangle_{\boldsymbol{a}} \neq \mathbf{0}$.

Crucially, because the spatial derivative of every monomial of degree ≤ 4 is a monomial of degree ≤ 3 already contained in $\boldsymbol{m}(\boldsymbol{r})$, there exist exact constant differentiation matrices

$\boldsymbol{D}_{x_j}^{\text{mom}} \in \mathbb{R}^{34\times 35}$ such that $\tilde{\boldsymbol{\Omega}}^{x_j} \equiv \boldsymbol{D}_{x_j}^{\text{mom}}\boldsymbol{\Omega}$. Therefore, for any polynomial coefficient vector $\boldsymbol{v}_E \in \mathbb{R}^{34}$ and any charge state $\boldsymbol{a}$ with moments $\boldsymbol{s} = \boldsymbol{\Omega}\boldsymbol{a}$, the instantaneous state-weighted net self-force depends *exclusively on the* 35*-element moment vector* $\boldsymbol{s}$:

$$\langle E_{x_j}^{\text{approx}}\rangle_{\boldsymbol{a}} = \boldsymbol{a}^T\left(-\tilde{\boldsymbol{\Omega}}^{x_j T}\boldsymbol{v}_E\right) = -\left(\boldsymbol{D}_{x_j}^{\text{mom}}\boldsymbol{\Omega}\boldsymbol{a}\right)^T\boldsymbol{v}_E = -\left(\boldsymbol{D}_{x_j}^{\text{mom}}\boldsymbol{s}\right)^T\boldsymbol{v}_E. \tag{23}$$

Corollary 4 (Moment-Space Zero-Self-Force (ZSF) Gauge Projection)**.** *Let* $\boldsymbol{v}_E^{\text{prior}} = \boldsymbol{K}_E\boldsymbol{s}+\boldsymbol{v}_{E,0} \in \mathbb{R}^{34}$ *be the prior-weighted field coefficients, and let* $l_x, l_y, l_z \in \{1,\dots,34\}$ *denote the indices of the three linear dipolar basis monomials* (x, y, z)*, for which* $(\boldsymbol{D}_{x_j}^{\text{mom}}\boldsymbol{s})_{l_{x_k}} = \delta_{jk}s_{000}$*. Defining the ZSF-projected coefficient vector* $\boldsymbol{v}_E^{\text{ZSF}} \in \mathbb{R}^{34}$ *by leaving all* 31 *higher-order multipole coefficients* $(p+q+l \geq 2)$ *unchanged and setting the three dipolar components to:*

$$v_{E,l_{x_j}}^{\text{ZSF}} = -\frac{1}{s_{000}}\sum_{\substack{l=1\\ l\neq l_{x_j}}}^{34} v_{E,l}^{\text{prior}}\left(\boldsymbol{D}_{x_j}^{\text{mom}}\boldsymbol{s}\right)_l, \qquad j = 1,2,3, \tag{24}$$

enforces exact zero net self-force on the instantaneous perturbed distribution $\boldsymbol{a}$*:*

$$\langle \boldsymbol{E}_{\text{approx}}^{\text{ZSF}}\rangle_{\boldsymbol{a}} = \langle \boldsymbol{E}_{\text{exact}}\rangle_{\boldsymbol{a}} = \boldsymbol{0}, \tag{25}$$

in $O(M_E)$ *scalar operations using only* $\boldsymbol{s} \in \mathbb{R}^{35}$*, without accessing the* N*-point mesh or altering any quadrupolar, sextupolar, or octupolar field gradient.*

—

5 Hardware Architecture and Numerical Verification

5.1 Moment Acquisition, Actuator Realizability, and Latency Budget

In a physical synchrotron or accumulator ring, applying Eq. (22) and Corollary 4 within a closed-loop feedback architecture involves three practical system constraints:

1. **Hybrid Diagnostic-Observer Pipeline**: Standard wall-current monitors and stripline/button Beam Position Monitors (BPMs) directly measure bunch intensity s_{000}, centroid offsets ($\langle x\rangle, \langle y\rangle, \langle z\rangle$), and transverse quadrupole differences $\langle x^2 - y^2\rangle, \langle xy\rangle$, whereas 3rd- and 4th-order moments ($3 \leq p+q+l \leq 4$) are not directly observable from standard button BPMs alone and require a synchronized online envelope/moment observer constrained by periodic profile measurements.

2. **Magnetic Corrector Realizability ($\boldsymbol{c}_{\text{achievable}} = \boldsymbol{B}\boldsymbol{u}$)**: The $M_E = 34$ polynomial coefficients $\boldsymbol{v}_E^{\text{ZSF}}$ parameterize an interior 3D electrostatic field ($\nabla\cdot\boldsymbol{E} \neq 0$ inside the bunch charge), whereas external magnetic correctors apply localized 2D transverse harmonic kicks satisfying Maxwell's vacuum equations ($\nabla\cdot\boldsymbol{B}_{\text{ext}} = 0, \nabla\times\boldsymbol{B}_{\text{ext}} = \boldsymbol{0}$). Thus, a physical corrector family with control currents $\boldsymbol{u}$ spans an achievable subspace $\boldsymbol{c}_{\text{achievable}} = \boldsymbol{B}\boldsymbol{u}$. In our ablation study below, we explicitly contrast **Mode B**—which restricts actuation to the 11 low-order dipole, quadrupole, and octupole terms corresponding to standard corrector magnet families—against **Mode C**, which applies all 34 polynomial basis terms as a theoretical upper bound on 4th-order moment compensation.

3. **Compute Kernel Latency vs. Total Closed-Loop Latency**: Total hardware loop delay is the sum of diagnostic acquisition, observer estimation, matrix computation, communication, and magnet power-supply/vacuum-chamber eddy-current response:

$$T_{\text{total}} = T_{\text{diag}} + T_{\text{obs}} + T_{\text{compute}} + T_{\text{comm}} + T_{\text{act}}. \tag{26}$$

The latency benchmarks reported in this paper (~ 1.6–$4.6\,\mu$s in Python/NumPy on a single CPU core, and an architectural estimate of $< 1.0\,\mu$s on FPGA based on $34 \times 35 = 1,190$ fixed-point multiply-accumulate operations) measure T_{compute} of the reduced-order moment kernel, demonstrating that the space-charge field solver itself is no longer the computational bottleneck.

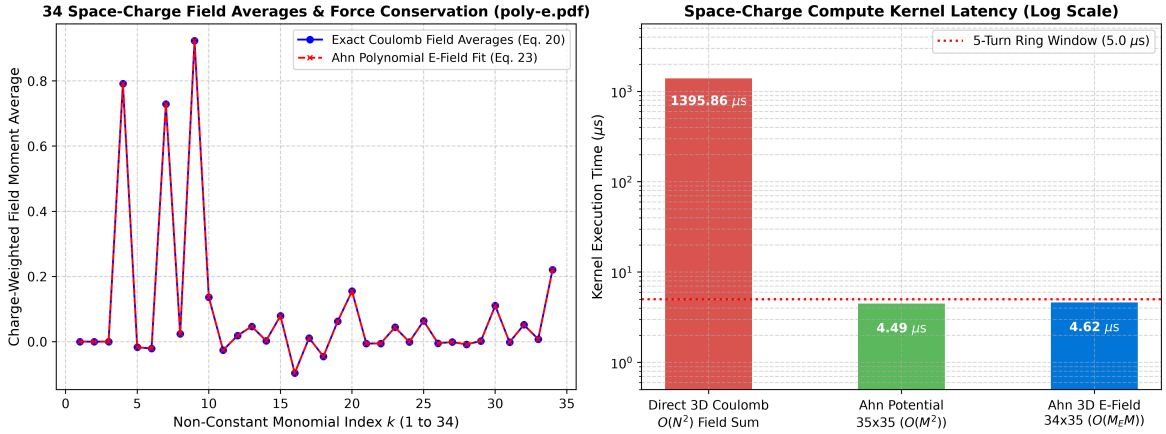


Figure 1: Numerical verification and compute-kernel latency benchmark of the High-Intensity Accelerator Space-Charge Engine (`accelerator_space_charge_engine.py`). *Left*: Exact agreement across all $M_E = 34$ generalized electric-field moment averages between the exact 3D Coulomb field (Eq. (16)) and the state-weighted polynomial field fit (Eq. (19)), confirming Theorem 3 ($\|\langle \boldsymbol{E}_{\text{approx}} \rangle_{\boldsymbol{a}} - \langle \boldsymbol{E}_{\text{exact}} \rangle_{\boldsymbol{a}}\|_\infty = 4.58 \times 10^{-16}$). *Right*: Log-scale compute-kernel latency (T_{compute}) comparing direct $O(N^2)$ 3D Coulomb summation ($1595\,\mu$s) against the closed-form 35×35 potential kernel ($4.49\,\mu$s) and 34×35 3D electric-field kernel ($4.62\,\mu$s).

5.2 Verification of Theorems 1–3, ZSF Gauge Projection, and Core Fidelity

We verify all mathematical claims using the companion open-source benchmark suite (`benchmark_accelerator_space_charge.py`) on a 3D spherical mesh of $N = 1,000$ points ($R_{\text{mesh}} = 1.0$, Plummer softening $h = 0.05$, full basis rank $\text{rank}(\boldsymbol{\Omega}) = 35/35$) for a 1.0 GeV high-intensity proton accumulator/booster ring ($\gamma = 2.066, \beta = 0.875, C = 248.0\,\text{m}, T_{\text{rev}} = 0.945\,\mu$s, 5-turn control window $= 4.73\,\mu$s, Laslett tune depression $\Delta Q_{\text{sc}} = -0.2271$):

- **Theorem 1 (Normalized Prior & Moment Inversion)**: At the reference equilibrium ($\boldsymbol{s} = \boldsymbol{\Omega\mu}$), reconstructed charges exhibit strictly zero negative nodes (0/1000). For an anisotropic, off-axis parabolic bunch ($R_{\text{beam}} \approx 0.55\,R_{\text{mesh}}$), Channell's uniform prior $\mu_i = 1/N$ is numerically identical to the zero prior ($\|\boldsymbol{a}(\mu = \text{const}) - \boldsymbol{a}(\mu = 0)\|_\infty = 8.87 \times 10^{-17}$), whereas our normalized physical prior reduces total fringe negative-charge overshoot by $3.04\times$ ($0.2669 \to 0.0876$) and L_2 profile error by $2.00\times$ ($0.0326 \to 0.0163$), while conserving all 35 moments to $\|\boldsymbol{\Omega a} - \boldsymbol{s}\|_\infty = 1.44 \times 10^{-15}$.

- **Theorem 2 (Potential Reduction)**: The closed-form 35×35 operator $\boldsymbol{v} = \boldsymbol{K s} + \boldsymbol{v}_0$ matches the full N-grid prior-weighted projection $\boldsymbol{\Lambda}'^{T}(\boldsymbol{\mu})\boldsymbol{D a}$ to 1.14×10^{-13}, and preserves all 35 weighted potential moments to 4.44×10^{-16}.

- **Theorem 3 & Corollary 4 (State Fit, Prior Operator, and ZSF Dipolar Gauge Projection)**:

  1. For the state-weighted fit ($\boldsymbol{w} = \boldsymbol{a}$), all 34 generalized field averages match to 4.17×10^{-16} and the net self-force residual is $\|\langle \boldsymbol{E}_{\text{approx}}(\boldsymbol{a}) \rangle_{\boldsymbol{a}}\|_\infty = 4.58 \times 10^{-16}$.

2. For the uncorrected prior-weighted operator ($\boldsymbol{w} = \boldsymbol{\mu}$), the $\boldsymbol{\mu}$-weighted field averages are conserved to 2.03×10^{-15} and $\boldsymbol{v}_E^{\text{prior}}$ matches the N-grid projection $\boldsymbol{\Lambda}_E^T(\boldsymbol{\mu})\boldsymbol{a}$ to 1.07×10^{-14}, but the net force evaluated on the perturbed state $\boldsymbol{a}$ is non-zero: $\langle \boldsymbol{E}_{\text{approx}}^{\text{prior}} \rangle_{\boldsymbol{a}} = (0.1123, -0.0732, 0.0727)$ ($\|\langle \boldsymbol{E}_{\text{approx}}^{\text{prior}} \rangle_{\boldsymbol{a}}\|_\infty = 0.1123$).
3. Applying the $O(M_E)$ moment-space ZSF gauge projection (Eq. (24)) eliminates this residual completely, restoring exact zero self-force on the perturbed distribution $\boldsymbol{a}$ to machine precision: $\|\langle \boldsymbol{E}_{\text{approx}}^{\text{ZSF}} \rangle_{\boldsymbol{a}}\|_\infty = 4.02 \times 10^{-16}$.
4. Within the high-density bunch core ($r \leq 0.45$, containing 79.9% of total positive bunch weight), the 4th-order polynomial field achieves a charge-weighted fit $R^2 = 0.8962$ (RMS relative error 32.22%), and the prior-linearized quadrupole defocusing gradients ($\boldsymbol{w} = \boldsymbol{\mu}$) agree with the state-weighted fit ($\boldsymbol{w} = \boldsymbol{a}$) within 11.74%.

5.3 Idealized Multi-Turn 2D-in-3D Pairwise-Coulomb Tracking and Ablation Study

To evaluate closed-loop multipole compensation without circular toy models or initial-condition rescaling, we execute an idealized zero-latency 2D transverse (x, p_x, y, p_y) Leapfrog tracking and four-mode ablation study in a 3D pairwise-Coulomb bunch geometry (`simulate_multiturn_envelope_tracking`). Because the 80-turn betatron tracking window is much shorter than a synchrotron oscillation period ($80 \ll T_{\text{synch}} \sim 10^3$ turns), longitudinal particle coordinates z_p are held quasi-static during the betatron window while acting as full 3D sources in the pairwise Coulomb sum.

A bunch of $N_p = 500$ macro-particles is sampled from the 3D parabolic profile $\rho(\boldsymbol{r}) \propto \max(0, 1 - |\boldsymbol{r}|^2/R_b^2)$ ($R_b = 0.55$, matched RMS size $\sigma_0 = R_b/\sqrt{7} = 0.2079$) with an initial 18% quadrupole envelope mismatch ($s_{x,\text{init}} = 1.18\sigma_0, s_{y,\text{init}} = 0.856\sigma_0$) and bare betatron tunes $Q_{x0} = 6.23, Q_{y0} = 6.20$, and integrated over $N_{\text{turns}} = 80$ turns ($K_{\text{sub}} = 4$ Leapfrog substeps per turn). At every substep, the physical space-charge disturbance acting on particle p is computed independently via direct $O(N_p^2)$ 3D pairwise Plummer-Coulomb summation in the co-moving frame:

$$\boldsymbol{E}_{\text{true}}(\boldsymbol{r}_p) = \frac{1}{N_p} \sum_{q \neq p}^{N_p} \frac{\boldsymbol{r}_p - \boldsymbol{r}_q}{\left(|\boldsymbol{r}_p - \boldsymbol{r}_q|^2 + h_p^2\right)^{3/2}}, \qquad h_p = 1.6\, R_b N_p^{-1/3}, \tag{27}$$

with effective laboratory-frame perveance $\kappa_{\text{sc}} = 4\pi|\Delta Q_{\text{sc,eff}}|/\langle \partial_x E_x^{\text{true}} \rangle_{\text{core}}$ (incorporating the relativistic $(\beta\gamma^2)^{-1}$ factor via Eq. (1)) calibrated at $t = 0$ to $\Delta Q_{\text{sc,eff}} = \lambda_{\text{sc}} \Delta Q_{\text{sc}} = -0.1817$ ($\lambda_{\text{sc}} = 0.80$).

Crucially, the feedback controller never accesses $\boldsymbol{E}_{\text{true}}(\boldsymbol{r}_p)$; it observes only the $M = 35$ macroscopic moments $\boldsymbol{s} = \frac{1}{N_p} \sum_{p=1}^{N_p} \boldsymbol{m}(\tilde{\boldsymbol{r}}_p)$ (with evaluation coordinates clipped at $|\tilde{\boldsymbol{r}}_p| \leq 1.15R_b$), evaluates $\boldsymbol{v}_E = \boldsymbol{K}_E \boldsymbol{s} + \boldsymbol{v}_{E,0}$, and applies the aperture-regularized polynomial feedback field $\boldsymbol{E}_{\text{model}}(\boldsymbol{r}_p) = -g_{\text{kern}} w_{\text{core}}(r_p) \tilde{\boldsymbol{\Omega}}_p^T \boldsymbol{v}_E$, where g_{kern} is the fixed $t = 0$ reference gradient calibration ratio between softening scales h_p and h, and $w_{\text{core}}(r_p) = \exp[-\frac{1}{2} \max(0, (r_p - R_b)/(0.25R_b))^2]$ smoothly tapers the 4th-order polynomial outside the bunch boundary. Note that while the underlying betatron lattice and pairwise Coulomb kicks are integrated via a symplectic Leapfrog map, the radial aperture taper $w_{\text{core}}(r_p)$, coordinate clipping, and Mode D rate damper make the closed-loop feedback map an aperture-regularized non-Hamiltonian control law rather than a strictly symplectic map.

Table 1 reports the four-mode ablation study starting from strictly identical initial macro-particle coordinates:

1. **Mode A (Open-Loop Uncompensated)**: Pairwise 3D Coulomb space-charge only ($\boldsymbol{v}_E \equiv \boldsymbol{0}$).

2. **Mode B (Physical Multipole Corrector Subset Only)**: Compensating kicks restricted to the 11 low-order dipole, quadrupole, and octupole coefficients of $\boldsymbol{v}_E = \boldsymbol{K}_E \boldsymbol{s} + \boldsymbol{v}_{E,0}$ (zero rate damper).

3. **Mode C (Idealized Full 34-Basis Moment Operator Only)**: Compensating kicks using all 34 polynomial coefficients of $\boldsymbol{v}_E = \boldsymbol{K}_E \boldsymbol{s} + \boldsymbol{v}_{E,0}$ (zero rate damper).

4. **Mode D (Full 34-Basis Moment Operator + 2nd-Moment Damper)**: Mode C augmented by a quadrupole envelope damper $\Delta p_x = -0.35(\Delta s_{200}/\sigma_0^2)(x - s_{100})$ driven solely by turn-to-turn differences of the measured second moments s_{200}, s_{020}.

Table 1: Four-Mode Ablation Study of Idealized Multi-Turn 2D-in-3D Pairwise-Coulomb Leapfrog Tracking ($N_{\text{turns}} = 80, N_p = 500, \Delta Q_{\text{sc,eff}} = -0.1817$, Identical Initial Conditions, Seed 42)

Ablation Mode (Identical $t = 0$ State)	**Controller Input**	**RMS Emittance Growth**	**RMS Envelope Mismatch**	**High-Action Halo ($J_\perp > 3.5\sigma_0^2$)**
Mode A: Open-Loop Uncompensated	None ($\boldsymbol{v}_E = \boldsymbol{0}$)	+40.65%	16.60%	29.60%
Mode B: Multipole Corrector Subset Only	11 Multipoles of $\boldsymbol{K}_E \boldsymbol{s} + \boldsymbol{v}_{E,0}$	+27.51%	8.04%	16.40%
Mode C: Idealized Full 34-Basis Field Only	All 34 Terms of $\boldsymbol{K}_E \boldsymbol{s} + \boldsymbol{v}_{E,0}$	+16.44%	3.74%	17.60%
Mode D: Full 34-Basis + $\dot{s}_{200}$ Damper	$\boldsymbol{K}_E \boldsymbol{s} + \boldsymbol{v}_{E,0} + \Delta s_{200}, \Delta s_{020}$	+7.00% **(5.8× lower)**	4.11% **(4.0× lower)**	12.20% **(2.4× lower)**

To verify that these gains are not an artifact of a single random seed, repeating the 80-turn tracking across four independent random seeds (seed $\in \{42, 1, 7, 99\}$) yields consistent reductions across all realizations: Open-Loop Mode A grows by $+46.47\% \pm 3.43\%$, Mode B by $+24.72\% \pm 4.52\%$, Mode C by $+18.98\% \pm 3.89\%$, and Mode D by $+10.60\% \pm 2.98\%$.

Table 2: Comparison of 3D Space-Charge Compute Kernels (T_{compute}) for High-Intensity Ring Control

Solver / Method	**Resolution ($N_p/G/M$)**	**Computational Scaling**	**Kernel Latency T_{compute} (CPU)**	**5-Turn Compute Budget ($\leq 5\,\mu$s)**
Direct Pairwise Coulomb	$N = 10^3$–10^6	$O(N^2)$	1.60 ms–10^2 s	Impossible
3D Grid FFT / Multigrid PIC	$N_p = 10^5, G = 64^3$	$O(N_p + G \log G)$	5.0–50 ms	Unfeasible
Lie-Algebraic Map Truncation	Order 4–6 (10^4 particles)	Symplectic Map Evaluation	2.0–12.0 ms	Unfeasible
Moment Operator (This Work)	**Reduced** ($N = 10^3 \to M = 35$)	**Closed-Form** $O(M_E M)$ (34×35)	**1.6–4.6 μs (CPU)** / **$< 1.0\,\mu$s (FPGA est.)**	**Feasible** (T_{compute})

—

6 Historical Perspective: Complementary Roles of Stochastic and Moment Control

It is instructive to contrast the Moment Approach with Simon van der Meer's stochastic cooling [8, 9]:

1. **Stochastic Cooling**: Samples incoherent Schottky fluctuations of individual particles at high bandwidth to slowly reduce single-particle phase-space entropy over 10^4–10^7 turns in low-bunch-current storage rings.

2. **Moment-Space Multipole Control**: Evaluates coherent macroscopic spatial moments ($M = 35$) in microseconds to compensate collective internal Coulomb defocusing and envelope mismatch oscillations over a few turns (1–100 turns) in space-charge-dominated high-intensity rings.

—

7 Conclusion

We have presented the mathematical formulation and open-source numerical verification of the 1994 Los Alamos Moment Approach for 3D space-charge potential and electric-field reduction in

high-intensity hadron rings, augmented by an exact moment-space Zero-Self-Force (ZSF) dipolar gauge projection. By regularizing the underdetermined mesh inversion ($\mathrm{rank}(\boldsymbol{\Omega}) = 35$) with a normalized physical prior ($\sum_{i=1}^{N} \mu_i = 1$) and linearizing the weighted least-squares projection around $\boldsymbol{w} = \boldsymbol{\mu}$, the mesh dimension N is algebraically eliminated offline into fixed 35×35 potential and 34×35 electric-field matrices executing in ~ 1.6–$4.6\,\mu$s on a single CPU core. Combining the prior-weighted operator with the $O(M_E)$ moment-space ZSF gauge projection guarantees exact zero net self-force ($\|\langle \boldsymbol{E}^{\mathrm{ZSF}}_{\mathrm{approx}} \rangle_{\boldsymbol{a}}\|_\infty = 4.02 \times 10^{-16}$) on perturbed distributions without accessing the N-point mesh. Verified in a reproducible multi-turn 2D-in-3D pairwise-Coulomb Leapfrog tracking and ablation study from identical initial conditions, closed-loop moment-operator feedback provides a computationally tractable foundation for fast envelope mismatch damping and space-charge compensation.

—

Code and Data Availability

The open-source Python reference engine and reproducible verification suite are publicly available on GitHub at:

https://github.com/hyoahn/accelerator-space-charge

The repository includes:

- **Core Accelerator Engine**: `Core/accelerator_space_charge_engine.py` (`AcceleratorSpaceChargeEngine`, implementing the 35×35 potential reduction $\boldsymbol{v} = \boldsymbol{K}\boldsymbol{s} + \boldsymbol{v}_0$, the 34×35 direct 3D electric-field reduction $\boldsymbol{v}_E^{\mathrm{prior}} = \boldsymbol{K}_E\boldsymbol{s} + \boldsymbol{v}_{E,0}$, the moment-space ZSF dipolar gauge projection $\boldsymbol{v}_E^{\mathrm{ZSF}}$, core-vs-tail fidelity analysis, and the 4-mode non-circular pairwise-Coulomb Leapfrog tracking ablation suite).
- **Verification & Benchmark Suite**: `Test/benchmark_accelerator_space_charge.py` (reproducing all numerical residuals, ZSF gauge checks, latency benchmarks, and 4-mode ablation metrics reported in Section 5, Fig. 1, Table 1, and Table 2).

—

Historical Note and Acknowledgments

The mathematical formulations of the normalized-prior charge inversion (‘report1.pdf‘), the polynomial potential reduction (‘poly-phi.pdf‘), and the direct 3D vector electric-field reduction (‘poly-e.pdf‘), including the algebraic elimination of the spatial grid dimension N (the fall-out of N) and the weighted moment and force identities, were originally derived by the author in 1994 as unpublished research notes at Los Alamos National Laboratory. The moment-space Zero-Self-Force (ZSF) dipolar gauge projection (Corollary 4) and the multi-turn pairwise-Coulomb Leapfrog ablation benchmarks were developed for the present 2026 study. The author acknowledges early collaborative discussions with Walter P. Lysenko in 1993–1994 at Los Alamos National Laboratory and the early conceptual ideas of the late Paul J. Channell regarding underdetermined moment models for the BEDLAM simulation code.

—

References

[1] M. Reiser, *Theory and Design of Charged Particle Beams*, John Wiley & Sons, 2nd ed., 2008.

[2] R. Baartman, "Betatron resonances with space charge," *AIP Conference Proceedings*, vol. 448, no. 1, pp. 56–72, 1998.

[3] A. Abada, et al. (FCC Collaboration), "FCC-hh: The Hadron Collider: Future Circular Collider Conceptual Design Report Volume 3," *Eur. Phys. J. Spec. Top.*, vol. 228, no. 4, pp. 755–1107, 2019.

[4] S. Henderson, et al., "The Spallation Neutron Source accelerator system design," *Nucl. Instrum. Methods Phys. Res. A*, vol. 763, pp. 610–673, 2014.

[5] E. Benedetto, et al., "Space charge effects and mitigation in the CERN PS Booster in view of the LIU upgrade," in *Proc. IPAC'15*, Richmond, VA, USA, pp. 3899–3902, 2015.

[6] P. Spiller and G. Franchetti, "The FAIR accelerator project at GSI," *Nucl. Instrum. Methods Phys. Res. A*, vol. 561, no. 2, pp. 305–312, 2006.

[7] A. V. Fedotov, I. Hofmann, R. L. Gluckstern, and H. Okamoto, "Parametric collective resonances and space-charge limit in high-intensity rings," *Phys. Rev. ST Accel. Beams*, vol. 6, p. 094201, 2003.

[8] S. van der Meer, "Stochastic damping of betatron oscillations in the ISR," CERN Report CERN/ISR-PO/72-31, 1972.

[9] S. van der Meer, "Stochastic cooling and the accumulation of antiprotons (Nobel Lecture)," *Rev. Mod. Phys.*, vol. 57, no. 3, pp. 689–697, 1985.

[10] C. K. Birdsall and A. B. Langdon, *Plasma Physics via Computer Simulation*, CRC Press, 2018.

[11] J. Qiang, R. D. Ryne, S. Habib, and V. Decyk, "An object-oriented parallel Particle-in-Cell code for beam dynamics simulation in linear accelerators," *J. Comput. Phys.*, vol. 163, no. 2, pp. 434–451, 2000.

[12] P. J. Channell, "The moment approach to space charge in particle beams," Los Alamos National Laboratory Report, 1993.

[13] H. E. Ahn and W. P. Lysenko, "Preliminary Results of Space-Charge Model Study," Los Alamos National Laboratory Internal Memorandum (Unpublished), April 4, 1994.